

AI DAILY

AI Daily 2026-07-21: AI Glasses Start Subtracting

MOONIX puts approximately 14.9 grams, about 16 hours, six microphones, AI capture and summary, and proactive AI into an everyday-wear loop.

The standard model emphasizes audio capture and AI questions while Pro adds a high-definition camera and audio-video recording; capability boundaries become an explicit product split.

Monako, RayNeo, Meta, and Android XR represent contrasting choices around agent HUDs, display, camera privacy, and platform trust.

- AI glasses
- wearability
- camera-free
- MOONIX
- coding agents
- Android XR
- privacy UX
- on-device AI

Sources: [Cover source](#)



confirmed product · official MOONIX page · China launch report · comparison with camera/display routes · The strongest product question today is not how many capabilities can fit into smart glasses. It is which capability earns the right to stay on the face. MOONIX removes display from the standard product, keeps audio, six microphones, capture, summary, and a 14.9-gram frame. Monako makes the opposite bet by putting a Linux coding-agent HUD on the face. Meta’s anti-tamper work shows that even after weight and utility, social trust remains a hardware state.

STRUCTURE BEFORE EVIDENCE

Read today through eight lanes: official products, reviews, community friction, wild products, research, patents, China, and global signals.

OFFICIAL

Official Desk

5 item(s)

REVIEWS

Review Desk

3 item(s)

COMMUNITY

Community Friction

1 item(s)

WILD

Wild Desk

3 item(s)

RESEARCH

Research Watch

1 item(s)

PATENT

Patent Watch

1 item(s)

CHINA

China / Global Compare

4 item(s)

GLOBAL

Global Signals

2 item(s)

OFFICIAL

Official Desk

DEVELOPER SURFACE · DEVELOPER SURFACE · OFFICIAL
ANDROID XR DOCS · PRE-LAUNCH HARDWARE

Android XR turns eyewear from a product
preview into a testable developer surface

CONFIRMED PRODUCT · CONFIRMED PRODUCT · OFFICIAL
PRIVACY CONTROL · REVIEW/COMMUNITY FRICTION

Meta turns capture-LED tamper detection into a
camera-level hardware constraint

CONFIRMED PRODUCT · CONFIRMED PRODUCT · OFFICIAL
LAUNCH · DEVELOPER SURFACE

Snap SPECS puts AI, spatial display, and agent
development into standalone glasses

DEVELOPER SURFACE · DEVELOPER SURFACE · ENTERPRISE
PLATFORM · NOT A SHIPPED CONSUMER DEVICE

Project Solara and MDEP turn agent-first
devices into an enterprise-management
problem

CONFIRMED PRODUCT · CONFIRMED PRODUCT · VOICE
INTERFACE · CHATGPT VOICE ROLLOUT

GPT-Live moves ChatGPT Voice from Q&A
entry point toward continuous conversation


Official

developer surface · developer surface · official Android XR docs · pre-launch hardware

2026-05-19 intelligent-eyewear preview · 2026-06-15 Developer Preview 4 · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

Android XR turns eyewear from a product preview into a testable developer surface

Product: Android XR intelligent-eyewear developer surface. Android XR intelligent-eyewear developer surface: At I/O 2026 Google split the Android XR eyewear route into audio glasses and display glasses and showed the direction with Samsung, Gentle Monster, Warby Parker, and Qualcomm. The concrete product evidence is Developer Preview 4: developers can use Android Studio emulator, Jetpack Projected, Compose Glimmer, and Gemini APIs to extend mobile apps to glasses. The hardware remains an autumn, select-market preview; the developer surface is public. This item is handled as developer surface; missing details remain source not stated.



WHAT IT IS
At I/O 2026 Google split the Android XR eyewear route into audio glasses and display glasses and showed the direction with Samsung, Gentle Monster, Warby Parker, and Qualcomm. The concrete product evidence is Developer Preview 4: developers can use Android Studio emulator, Jetpack Projected, Compose Glimmer, and Gemini APIs to extend mobile apps to glasses. The hardware remains an autumn, select-market preview; the developer surface is public.

SPECS / STACK
The public stack includes Android XR SDK Developer Preview 4, Android Studio emulator, Jetpack Projected, Jetpack Compose Glimmer, Device Availability API, Gemini Live API, Firebase AI Logic, Kotlin, Unity, Unreal, Godot, ARCore Geospatial API, and OpenXR. Google describes headsets, wired XR glasses, audio glasses, and display glasses as distinct form factors. Samsung/Google intelligent eyewear is scheduled for select markets in fall 2026. Lens parameters, cameras, processors, weight, battery, price, exact Android version, and edge/cloud model split are source not stated.

PAIN POINTS
The surface targets three early glasses problems: teams do not know where to start, phone and eyewear state drift apart, and transparent displays cannot reuse phone layouts directly. Projected makes cross-device projection a library; Device Availability brings ‘is the eyewear being worn?’ into lifecycle logic; Compose Glimmer supplies readability and touchpad components. For users the goal is fewer phone pulls. For developers the goal is lowering the migration cost from demo to maintainable app.

NEW TECH
The move is not simply that Gemini can answer questions. It is that the system exposes agent, input, and device state to applications. Apps can use Gemini Live API and function calling for real-time voice, while Android Intelligence and AppFunctions let an agent discover and execute declared services, data, and actions. Projected and Glimmer define a field-of-view constraint: short, scannable, and operable by touchpad or voice. Device Availability makes wearable context an explicit state developers must handle.

LIMITS / UNKNOWNNS
The sources do not state a unified glasses specification, measured optical-display legibility, voice error rates in noise, Gemini latency or cost, offline behavior after phone disconnect, or how AppFunction permissions appear on glasses. Notification density, walking attention, and bystander understanding of audio and display require hardware testing. A Developer Preview proves an API exists, not that daily wear works.

HOW IT WORKS
A wearer can ask Gemini for directions, a nearby coffee shop, pickup ordering, notification summaries, calendar events, photos, and live translation. Display glasses can translate menus and signs in view, while audio glasses return spoken help. A developer starts with an Android app, uses Jetpack Projected to bring a companion experience to glasses, uses Device Availability to adapt lifecycle behavior when the device is worn, and uses Compose Glimmer for text and touchpad components on optical see-through displays. Gemini Live API and function calling provide real-time voice and action entry points.

USE CASES
Official examples include asking for walking directions, finding a coffee shop and ordering pickup on the route, receiving summarized messages, adding calendar events, translating speech, translating a menu or sign in view, and taking a photo without the phone. For developers, the job is turning a mobile app into a glanceable companion: return to the phone when glasses are unavailable and compress complex tasks into voice, short text, or touchpad controls while walking. The emulator lets teams test flows and state before hardware access.

USER VOICE
Source not stated.

AVAILABILITY
Developer Preview 4, Android XR documentation, Gemini guidance, and the emulator are public. The Catalyst Program offers accepted developers pre-release hardware, support forums, and launch guidance. Google and Samsung describe intelligent eyewear as coming to select markets in fall 2026; model, country list, price, preorder, prescription options, and shipping date are source not stated. Development and simulation are available; consumer availability is not confirmed.

PRODUCT READ
The most important Android XR output is not an unpriced pair of glasses. It is a set of interfaces that can move a mobile app into the user’s sight, ear, and agent loop. Projected, Glimmer, and Device Availability turn hardware constraints into a developer contract. Ecosystem success depends on stable autumn hardware, explainable permissions, and developers compressing complex apps into legible, stoppable, recoverable micro-flows.


Official

confirmed product · confirmed product · official privacy control · review/community friction

2026-07-16 official Meta FAQ · 2026-07-17 Android Central trust interview · 2026-07-20 follow-up

Meta turns capture-LED tamper detection into a camera-level hardware constraint

Product: Meta AI Glasses. Meta AI Glasses: Meta AI Glasses are everyday glasses with cameras, microphones, speakers, and Meta AI. The product change worth recording today comes from Meta’ s July FAQ: the capture LED becomes an external signal of camera activity, and from the second generation onward the camera is disabled when the system detects that the LED has been covered, disabled, physically modified, or destroyed. The product question moves from whether privacy is promised to what bystanders can see and what the device will allow. This item is handled as confirmed product; missing details remain source not stated.



Meta official FAQ screenshot; capture LED, camera disabling, device-side import, and bystander privacy follow the source.

WHAT IT IS
Meta AI Glasses are everyday glasses with cameras, microphones, speakers, and Meta AI. The product change worth recording today comes from Meta’ s July FAQ: the capture LED becomes an external signal of camera activity, and from the second generation onward the camera is disabled when the system detects that the LED has been covered, disabled, physically modified, or destroyed. The product question moves from whether privacy is promised to what bystanders can see and what the device will allow.

SPECS / STACK
The source confirms a capture LED, camera, microphones, speakers, local glasses storage, phone import, Meta AI, and a companion app. The FAQ does not state processor, RAM, camera model, recording bitrate, edge/cloud model split, or LED brightness and detection thresholds, so those details remain source not stated. Second-generation camera disabling and the tamper-detection behavior are public constraints; they should not be expanded into an absolute privacy guarantee.

PAIN POINTS
The system addresses two problems: the wearer wants an always-available camera and assistant, while bystanders do not want invisible recording. Binding the LED to camera permission reduces the ambiguous state in which a user covers the light yet keeps recording. Local capture storage and deliberate import reduce the intuitive risk of everything going to the cloud automatically. The cost is that hardware and software must judge tampering, and false positives can remove camera access.

NEW TECH
The product move is to encode a social norm in a hardware state machine: the capture LED is part of camera permission, not merely an indicator. Meta also says it is improving detection of physical modification and acting against services that sell LED-tampering methods. The design pattern is visible state, capability coupling, and recovery. Privacy moves out of a settings page into a physical feedback layer shared by wearer and bystanders.

LIMITS / UNKNOWNNS
The sources do not state false-positive rate, detection latency, recovery after natural obstruction, whether an AI query can run without camera access, or whether all sensors and cloud processing stop when the camera is disabled. Independent testing has not established that bystanders can see the LED in daylight, crowds, or at distance. A camera constraint is not the same thing as a complete data-processing boundary.

HOW IT WORKS
The wearer uses voice, temple controls, or the companion app to take photos and video, listen to audio, ask Meta AI questions, and decide when to import captures to the phone gallery. The LED blinks briefly for a photo and continuously during video; the wearer hears a shutter sound. Meta says gallery captures are stored on the glasses until the user imports them to the phone. If the LED is blocked or tampering is detected, the capture path is blocked; recovery requires the LED to be unblocked and no longer in a tampered state.

USE CASES
Documented uses include hands-free photos and video, music and podcasts, questions to the assistant, asking about a live scene, and importing selected captures for sharing. The wearer gets a lower-friction camera and assistant entry point; people nearby get a visible recording cue. Real settings include gatherings, public transit, retail, and workplaces, where acceptance depends on whether bystanders understand the LED and whether the wearer can explain the recording behavior.

USER VOICE
Source not stated.

AVAILABILITY
Meta’ s FAQ is public and the AI Glasses product family remains on sale, with details varying by region and generation. The FAQ does not provide a new universal price, shipping date, firmware version, model-by-model coverage, or every regional control path; those remain source not stated. Android Central reports an early-July software update that disables the camera when LED tampering is detected, while this dossier prioritizes Meta’ s own behavior description.

PRODUCT READ
Meta turns a privacy dispute into a testable product action: if the light is covered or damaged, the camera stops. That is closer to real interaction than a policy statement, while independent testing, clear status feedback, and local rules are still needed. The next contest for AI glasses includes whether wearer and bystander can share the same recording state, rather than leaving the boundary visible only inside an app.

Official confirmed product · confirmed product · official launch · developer surface

2026-06-16 official launch · 2026-07 current pre-order watch · 2026-07-14 follow-up · 2026-07-14 follow-up · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

Snap SPECS puts AI, spatial display, and agent development into standalone glasses

Product: Snap SPECS. Snap SPECS: Snap SPECS is a standalone see-through AR glasses computer positioned between conventional AI glasses and an enclosed headset. Snap describes it as a wearable computer bringing AI assistance, work tools, entertainment, and shared experiences into the physical world without a puck or tether. The preorder announcement, specifications, and developer tools confirm the product surface; ecosystem scale and daily retention remain open. This item is handled as confirmed product; missing details remain source not stated.



Snap official SPECS launch visual; price, weight, display, battery, and APIs follow the release.

WHAT IT IS

Snap SPECS is a standalone see-through AR glasses computer positioned between conventional AI glasses and an enclosed headset. Snap describes it as a wearable computer bringing AI assistance, work tools, entertainment, and shared experiences into the physical world without a puck or tether. The preorder announcement, specifications, and developer tools confirm the product surface; ecosystem scale and daily retention remain open.

SPECS / STACK

Snap lists a 132-gram 47 mm model, a 136-gram 52 mm model, a 51-degree field of view, 16 million colors, two Snapdragon processors, 7 ms motion-to-photon latency, and up to four hours of mixed-use battery. The case supplies four additional charges, for up to 20 mixed-use hours. The system includes display, waveguide, electrochromic lenses, hand tracking, Snap OS, Lens Studio, and AI assistance. Processor model, RAM, camera details, networking, and edge/cloud split are source not stated.

PAIN POINTS

SPECS targets the conflict between phones that make people look down, headsets that isolate them, and AI glasses with limited displays. A larger private display and standalone compute reduce phone and tether dependence; Snap OS and tooling reduce the hardware-to-app gap. New costs include visual-attention competition, gesture fatigue, public capture boundaries, and explaining when an AI system sees, processes, or stores information.

NEW TECH

The increment is a combination of optics and waveguide, dual Snapdragon processors, hand tracking, Snap OS, a Spatial Benchmark, a Migration Agent, and agentic development in Lens Studio. Snap also says the glasses ask before sensitive access, show an LED during recording, prioritize on-device processing, and give users control over storage, sync, sharing, and deletion. The shift is a buildable and evaluable spatial-agent surface, not simply a chatbot inside glasses.

LIMITS / UNKNOWNNS

Open questions include whether four hours of mixed use survives real spatial-AI workloads, whether 132–136 grams is comfortable for long wear, which complete input path the 7 ms metric covers, and how Lenses obtain consent, explain data use, and recover from errors. The price places SPECS in an early-adopter bracket. A launch demo is not long-term usability evidence.

HOW IT WORKS

A wearer uses spatial display, hand tracking, voice, and Snap Lenses. Official examples include directions, spatial measurement, contextual AI, screen casting, whiteboards, and educational overlays on real objects. Developers build, debug, and publish Lenses in Lens Studio. A Migration Agent, Spatial Benchmark, Native Development Kit, and agentic-development preview extend the entry point from filters to spatial agents.

USE CASES

Confirmed use cases include directions and spatial measurement, bringing a screen or whiteboard into a work setting, educational and creative Lenses, contextual AI prompts, and migrating projects to SPECS. For workers, the value is keeping body and environment connected; for creators, it is spatial software rather than a phone filter. Weight, battery, and developer content will determine whether it becomes daily wear.

USER VOICE

Source not stated.

AVAILABILITY

Snap says SPECS is available for preorder at \$2,195 with a refundable \$200 deposit and is expected to ship in fall 2026 in the United States, United Kingdom, and France. Developer previews are rolling out in Claude Code, Codex, and Cursor. Consumer feature scope, prescription options, regional expansion, subscription, and third-party agent permissions are source not stated.

PRODUCT READ

SPECS is the clearest agent-first hardware proposition in this issue because spatial display, standalone computing, and developer tooling arrive together. Snap must prove that people will wear a \$2,195, four-hour computer and that developers can build agents that are legible, stoppable, and shareable. If the ecosystem forms, AI glasses move from ‘does it have a camera?’ to ‘who controls the spatial software layer?’

Official

developer surface · developer surface · enterprise platform · not a shipped consumer device

2026-06-02 official Microsoft Build post · 2026-07 platform watch · 2026-07-14 follow-up · 2026-07-14 follow-up · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

Project Solara and MDEP turn agent-first devices into an enterprise-management problem

Product: Microsoft Project Solara / MDEP. Microsoft Project Solara / MDEP: Project Solara is Microsoft’ s Build 2026 chip-to-cloud platform for agent-first enterprise devices; MDEP is the operating and management foundation around it. Microsoft places it across meeting-room systems, enterprise phones, digital signage, industrial and IoT endpoints, and AI-enabled edge devices. It is not a shipped terminal. It moves the question of how devices host agents from individual apps into identity, management, and security infrastructure. This item is handled as developer surface; missing details remain source not stated.

The request is blocked.

20260713T140516Z-16c8b64d88@pmhC10PVGud800000000pg00000000aevg

Microsoft MDEP post screenshot; Project Solara, Entra, Intune, Defender, and fleet management follow the post.

WHAT IT IS Project Solara is Microsoft’ s Build 2026 chip-to-cloud platform for agent-first enterprise devices; MDEP is the operating and management foundation around it. Microsoft places it across meeting-room systems, enterprise phones, digital signage, industrial and IoT endpoints, and AI-enabled edge devices. It is not a shipped terminal. It moves the question of how devices host agents from individual apps into identity, management, and security infrastructure.
SPECS / STACK Microsoft confirms an enterprise-grade OS, Entra identity, Intune management, Defender for Endpoint protection, a common fleet foundation, and Project Solara’ s chip-to-cloud direction. Media coverage adds a lightweight edge OS, Azure agent services, persistent cloud state, an agent dispatcher, and a task manager, reportedly based on AOSP; those details need later official documentation. Chip, AOSP branch, API, SDK, device models, price, shipping, and public developer eligibility are source not stated.
PAIN POINTS Traditional enterprise devices split work by app and model; agent-first devices will send tasks across terminals, services, and agents. MDEP tries to pull identity, management, threat protection, lifecycle, and trust into one platform. Users may get fewer sign-ins and clearer boundaries. The risk is that when an agent acts for an organization, responsibility, permissions, logs, and human takeover must be more explicit than in the app era.
NEW TECH The pattern combines an agent dispatcher, task manager, cloud state, and enterprise OS controls, letting devices inherit Entra, Intune, and Defender. If the AOSP edge-OS and Azure-agent route holds, a device becomes an embodied endpoint rather than a container for Office or Android apps. The key test is cross-device state and administrator visibility: can managers see what the agent did, not merely that a device is online?
LIMITS / UNKNOWNNS Open questions include what OS and API MDEP exposes, which Solara components are available, whether the task manager is visible to administrators, how errors roll back, how cross-device identity is least-privileged, and how AOSP coexists with Windows and Android. Enterprise buyers also need residency, audit, compliance, offline behavior, and an exit path; current sources provide no empirical details.


HOW IT WORKS An end user works in a meeting room, factory, clinic, or front desk through an agent-driven device. The agent is dispatched according to task and context, so the user experiences a work goal, prompt, or device action instead of launching a traditional app. Administrators define scope, policy, and lifecycle through Entra, Intune, and Defender for Endpoint. The promise is that agents inherit those guarantees rather than rebuilding identity, device management, and threat protection in every application.
USE CASES The enterprise scope covers meeting-room systems, industrial endpoints, clinics, front desks, digital signage, enterprise phones, and other edge devices. A task can be routed to an agent while the platform activates the appropriate agent according to context and administrators manage identity, policy, and security centrally. This addresses the need to avoid rebuilding accounts, deployment, updates, and audit controls for every new form factor; partner hardware and pilots still have to validate the experience.
USER VOICE Source not stated.
AVAILABILITY The MDEP post was published on June 2, 2026 and describes Project Solara as an announced platform, with no consumer purchase path or end-device release date. Enterprise partners, SDK access, developer enrollment, Azure regions, license pricing, and public pilots are source not stated. This remains a developer-surface and enterprise-platform watch, not a confirmed shipped product.
PRODUCT READ Project Solara and MDEP place agent UX back in the system foundation. A device becomes an execution node constrained by identity, policy, lifecycle, and security. Microsoft has not delivered a normal-user product to test, so the claim should remain measured. For enterprise HCI, the signal is clear: design must connect what the agent did, for whom, on which device, and who can stop it to the management plane.

Official

confirmed product · confirmed product · voice interface · ChatGPT Voice rollout

2026-07-08 official release · 2026-07-11 follow-up · 2026-07-12 follow-up · 2026-07-13 follow-up · 2026-07-14 follow-up · 2026-07-14 follow-up · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

GPT-Live moves ChatGPT Voice from Q&A entry point toward continuous conversation
Product: OpenAI GPT-Live. OpenAI GPT-Live: GPT-Live is OpenAI’s new generation of voice models released on July 8, 2026, positioned as the model powering ChatGPT Voice. This is a product-interface release, not merely a text-model announcement. The user-facing change is that voice interaction is meant to feel more like a continuous conversation: the assistant can maintain conversational flow while heavier work is delegated to a frontier model in the background. OpenAI says that at launch GPT-Live uses GPT-5.5 in the background and will be updated as new frontier models are released.. This item is handled as confirmed product; specs, price, regions and dates use cited source claims only, and missing details remain source not stated.



Source-backed screenshot from OpenAI GPT-Live official release page for the new ChatGPT Voice model.

WHAT IT IS
GPT-Live is OpenAI’s new generation of voice models released on July 8, 2026, positioned as the model powering ChatGPT Voice. This is a product-interface release, not merely a text-model announcement. The user-facing change is that voice interaction is meant to feel more like a continuous conversation: the assistant can maintain conversational flow while heavier work is delegated to a frontier model in the background. OpenAI says that at launch GPT-Live uses GPT-5.5 in the background and will be updated as new frontier models are released.

SPECS / STACK
The cited official sources support GPT-Live as a new voice model for ChatGPT Voice, delegation of web search, deeper reasoning, and complex work to a frontier model, launch-time use of GPT-5.5 in the background, and a July 8, 2026 OpenAI livestream replay. Media reports add GPT-Live-1, Mini, paid and free access tiers, global rollout, and live translation, but this issue treats those as supplements unless directly visible in OpenAI’s product page. Audio sampling, latency numbers, full language list, pricing, API naming, enterprise controls, and retention settings are source not stated.

PAIN POINTS
GPT-Live attacks two voice-AI frictions. First, older assistants often interrupt at pauses, forcing users to speak in machine-shaped rhythms. Second, complex work creates silent waiting, leaving the user unsure whether the system understood or is still working. GPT-Live combines turn-taking, background delegation, and short status cues so waiting becomes a legible conversational state. The risk is also clear: if delegation is slow, explanations are vague, or translation is wrong, voice users have less review and rollback space than text users.

NEW TECH
The product novelty is orchestration inside the voice interface. A foreground voice model handles natural turn-taking and presence, while a background frontier model handles heavier reasoning, search, or complex work. That split matters for future voice agents: the system needs to both stay with the user and complete tasks, while expressing background state in short, natural, low-interruption cues.

LIMITS / UNKNOWNNS
Open questions include real latency, weak-network behavior, noisy-room recognition, multi-speaker handling, privacy and audio retention, live-translation quality, developer API support, admin controls, and whether users can see when the background model changes. Voice also lacks the review buffer of text, so correction and undo need to be much closer to the interaction.

HOW IT WORKS
The user enters ChatGPT Voice and speaks naturally instead of treating the system as a strict record-one-turn, wait-one-turn interface. GPT-Live is meant to handle pauses, follow-up thoughts, and conversational flow. When a question requires web search, deeper reasoning, or more complex work, it can acknowledge the work with a short conversational cue, delegate the hard part to the background model, and return the result into the ongoing voice exchange. Media coverage also describes live translation, fewer interruptions, and instructions such as asking the assistant to stay quiet until addressed, but availability and account rules should be checked against official product behavior.

USE CASES
The concrete jobs are hands-free Q&A while walking or driving, translation, long-task planning, meeting preparation, spoken brainstorming, language learning, and information lookup when a screen is inconvenient. The stronger product move is from answering one spoken question to accompanying a user through work. A user can speak through a goal while the system keeps conversational presence and sends heavier retrieval or generation into the background. In HCI terms, voice becomes both input and status feedback, not only dictation.

USER VOICE
The evidence today is official release material plus media and developer commentary, not long-term retention data. TechRadar and MacRumors focus on naturalness, fewer interruptions, real-time translation, and replacement of the existing ChatGPT Voice experience. Ordinary-user evidence about accents, noisy spaces, privacy prompts, correction loops, multilingual accuracy, and enterprise controls still needs observation; source not stated.

AVAILABILITY
OpenAI’s official page says GPT-Live has launched and powers ChatGPT Voice. Specific account tiers, countries, client versions, workspace policy controls, API availability, and rollout timing are source not stated unless explicitly provided by the cited sources.

PRODUCT READ
GPT-Live is the cover story because it pushes AI product entry points from the chat box toward continuous spoken collaboration. The practical product judgment is that voice agents win on turn-taking, waiting feedback, correction, and privacy cues as much as model intelligence. If those states are clear, voice becomes a natural shell for an AI operating environment. If not, it remains a compelling demo wrapped around a fragile Q&A loop.

REVIEWS

Review Desk

CONFIRMED PRODUCT · CONFIRMED PRODUCT · OFFICIAL
PRODUCT PAGE · HANDS-ON REVIEW FRICTION

RayNeo X3 Pro makes display AI glasses
buyable, but not yet easy to use

CONFIRMED PRODUCT · CONFIRMED PRODUCT · OFFICIAL
PRODUCT · REVIEW/COMMUNITY FRICTION

HTC VIVE Eagle puts camera, translation, and
local data into lifestyle eyewear

CONFIRMED PRODUCT · CONFIRMED PRODUCT ·
REVIEW/COMMUNITY FRICTION

Even G2 trades cameras and speakers for
productivity and social acceptability

Reviews

confirmed product · confirmed product · official product page · hands-on review friction

2026-07-14 TechRadar review · current RayNeo UK product page

RayNeo X3 Pro makes display AI glasses buyable, but not yet easy to use

Product: RayNeo X3 Pro AI+AR Glasses. RayNeo X3 Pro AI+AR Glasses: RayNeo X3 Pro is a standalone, purchasable AI+AR glasses product that puts Gemini Live, RayNeo AIOS, full-colour MicroLED, real-time translation, navigation, five-way touch, and Creator Mode into one device. A July 14 TechRadar review brings the product back from launch language to daily use: the binocular display, outdoor legibility, and spatial positioning are strong, while 76-gram weight, battery life, and app-ecosystem friction still make it feel like a first-generation computer. This item is handled as confirmed product; missing details remain source not stated.

WHAT IT IS
RayNeo X3 Pro is a standalone, purchasable AI+AR glasses product that puts Gemini Live, RayNeo AIOS, full-colour MicroLED, real-time translation, navigation, five-way touch, and Creator Mode into one device. A July 14 TechRadar review brings the product back from launch language to daily use: the binocular display, outdoor legibility, and spatial positioning are strong, while 76-gram weight, battery life, and app-ecosystem friction still make it feel like a first-generation computer.

SPECS / STACK
RayNeo’ s product page lists a 640 by 480 full-colour MicroLED display, 6,000-nit peak brightness, Snapdragon AR1, 12MP camera, Gemini Live, Bluetooth 5.3, Wi-Fi 6, and Creator Mode. TechRadar’ s specification table adds 4GB LPDDR5 RAM, 32GB storage, 30-degree field of view, 60Hz, Android-based RayNeo AIOS, Google Gemini 2.5 Beta, a Sony IMX681 12MP colour camera plus a monochrome positioning camera, 6DoF and SLAM, 76 grams, a 245mAh battery rated at roughly one to five hours depending on use, USB-C charging in about 38 to 45 minutes, and translation across 14 languages with an approximately 2.1-second response. Edge/cloud split, recording bitrate, and regional API permissions are source not stated.

PAIN POINTS
The product addresses the limitation of camera-only glasses that cannot put information into the wearer’ s view, reducing phone pulls for navigation, translation, and notifications. Binocular MicroLED and 6,000 nits make outdoor legibility a measurable claim, while spatial positioning keeps overlays from drifting. The costs are equally concrete: 76 grams add pressure to nose and temples; mixed display, camera, and AI use may reduce the battery to an hour or a few hours; and the official app ecosystem still requires technical workarounds in some flows.

NEW TECH
The technical combination is binocular full-colour MicroLED through waveguides, RayNeo’ s Firefly Optical Engine, Snapdragon AR1, 6DoF/SLAM, AIOS, Gemini Live, and Creator Mode. The important interaction change is not another chatbot. It is making model output a spatial object: text must be stable, short, and dismissible; spatial state must recover; battery must become part of task design. The device moves display AI glasses from ‘can display’ to ‘is the display worth wearing continuously?’

LIMITS / UNKNOWNNS
The review gives the product a 3/5 first-generation verdict. Unknowns include mixed-load continuous runtime with display, camera, and AI active; Gemini consistency across languages and regions; whether the app ecosystem can escape workarounds; the battery cost of outdoor brightness; long-term comfort at 76 grams; and how open Creator Mode really is. Official specifications do not replace independent tests of charging frequency, heat, recognition errors, privacy cues, and offline fallback.


HOW IT WORKS
The wearer uses Hey RayNeo or the five-way touch panel on the right temple to start questions, translation, navigation, notifications, and capture. Binocular display places text and directions in the central field of view; open-ear speakers return audio; phone and RayNeo AIOS supply account, network, and application support. The review describes 6DoF plus SLAM and Falcon Image spatial positioning, but the real flow is constrained by app availability, region, model defaults, and remaining battery. The user reads in the field of view while deciding when to hand the task back to the phone.

USE CASES
Concrete jobs include outdoor navigation, street and menu translation, hands-free capture, notification handling, visual questions, and anchoring AR information to the physical scene. Binocular display is useful for routes, short text, and status cues without repeated phone pulls; Creator Mode opens a content and developer surface. The review also makes clear that this is not a general-purpose movie screen. Its value depends on a rhythm of glance briefly, act once, and return attention to the environment.

USER VOICE
Source not stated.

AVAILABILITY
The RayNeo UK product page lists £1,169 on sale from £1,299, with a 30-day price match, 30-day return and exchange, one-year warranty, and an estimated two-to-four-day delivery window. TechRadar reports a US price of \$1,169, an early-bird price of \$1,099, and a standard retail price of \$1,299. Prescription inserts are available through Lensology from about \$49/£49. Country inventory, software version, support policy, and Gemini feature differences vary by local page; missing details remain source not stated.

PRODUCT READ
RayNeo X3 Pro proves that display AI glasses can be bought as a standalone product: binocular MicroLED, spatial positioning, and outdoor legibility deliver capabilities camera-only glasses cannot. It also exposes first-generation constraints with unusual clarity. Weight, battery, price, and ecosystem decide whether ‘seeing the future’ becomes a daily task surface. Treat it as a technically accomplished confirmed product whose everyday completion is still incomplete.



TechRadar review page dated July 14, 2026: binocular MicroLED strengths and battery, weight, and ecosystem friction in one source.

Reviews

confirmed product · confirmed product · official product · review/community friction

2025-08-14 official launch · 2026-07-01 HardwareZone review · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

HTC VIVE Eagle puts camera, translation, and local data into lifestyle eyewear

Product: HTC VIVE Eagle. **HTC VIVE Eagle:** VIVE Eagle is HTC’ s camera-based AI glasses system, paired with the VIVE Connect mobile app and a VIVE AI cloud service. It puts voice assistance, hands-free photos, music, reminders, notes, restaurant recommendations, and photo translation into an everyday frame. The form factor is closer to audio glasses with visual input than to an immersive display product, aimed at travel and fast capture. This item is handled as confirmed product; missing details remain source not stated.

Access Denied

You don't have permission to access "http://www.vive.com/us/product/vive-eagle/overview/" on this server.
Reference #18.9329c817.1784037857.76935941
https://errors.edgesuite.net/18.9329c817.1784037857.76935941

HTC VIVE official product-page screenshot; 12MP camera, VIVE AI, 13-language translation, battery, and privacy follow the source.

WHAT IT IS

VIVE Eagle is HTC’ s camera-based AI glasses system, paired with the VIVE Connect mobile app and a VIVE AI cloud service. It puts voice assistance, hands-free photos, music, reminders, notes, restaurant recommendations, and photo translation into an everyday frame. The form factor is closer to audio glasses with visual input than to an immersive display product, aimed at travel and fast capture.

SPECS / STACK

HTC lists a 12MP ultra-wide camera, a frame under 49 grams, a 235mAh battery, up to 36 hours of standby, about 4.5 hours of continuous music playback, roughly 50% charge after ten minutes, open-ear audio, ZEISS sun lenses, a recording LED, AES-256 encryption, and translation across 13 languages. The system consists of the glasses, VIVE Connect, and a cloud AI service platform. Chip, RAM, recording resolution, connectivity method, and the edge/cloud split are source not stated.

PAIN POINTS

VIVE Eagle targets the action of seeing something and then reaching for a phone by putting camera, voice, and translation behind one entry point. Local storage and anonymization claims address concern about photos and voice being used for training. The review evidence shows that pairing, app setup, network conditions, and answer quality still shape the experience. The flow is shorter, while third-party models, cloud services, and data boundaries become part of the glasses product.

NEW TECH

The combination is a 12MP ultra-wide camera, open-ear audio, multi-model VIVE AI access, 13-language photo translation, local glasses storage, and AES-256. HTC also makes recording stop when the LED is obstructed or the glasses are removed. The design value depends on feedback: users need to know whether a photo was captured, whether translation started, which model supplied the answer, and what happens after a network failure.

LIMITS / UNKNOWNNS

Claims about local storage, no training use, and anonymization still need technical audit and independent testing; third-party model requests change the data path. Public sources do not state capture upload latency, translation accuracy, measured audio leakage, continuous AI runtime, prescription availability by region, or a phone-disconnected fallback mode. Review evidence remains cautious about accuracy, connectivity, and everyday value.

HOW IT WORKS

The wearer uses a phrase such as ‘Hey VIVE’ or a shortcut to take a photo, record, create a reminder, write a note, ask for a recommendation, play music, or translate. The camera capture is sent to VIVE AI and translation comes back through open-ear audio; the phone app manages content and settings. HTC supports third-party model platforms including OpenAI GPT and Google Gemini, so model availability, network access, and response quality are part of the real interaction flow.

USE CASES

Documented jobs include photographing a street scene and translating a menu while traveling, recording reminders by voice, listening to music and taking calls outdoors, finding a restaurant, translating what the camera sees into speech, and creating first-person records. The product targets users who want free hands, environmental awareness, and fewer phone pulls. Open-ear audio preserves ambient sound while the camera makes translation and visual understanding more concrete than an audio-only device.

USER VOICE

Source not stated.

AVAILABILITY

HTC announced an initial Taiwan launch at NT\$15,600 in Berry, Coffee, Grey, and Black. The US product page is public, while actual retail regions and delivery should be checked locally. Launch material mentions two years of VIVE AI Plus, but subscription coverage and regional eligibility are source not stated. HardwareZone’ s July 2026 review supplies hands-on friction and does not replace HTC’ s specification pages.

PRODUCT READ

VIVE Eagle is a complete, purchasable camera-AI glasses route: voice triggers, the camera sees, open-ear audio returns the answer, and phone plus cloud services fill the gaps. Its strength is concrete use cases; its risk is dependence on connectivity, third-party models, and clear data explanations. The key product interaction remains: what is being captured, who is processing it, and when does it stop?

Reviews

confirmed product · confirmed product · review/community friction

2026-07-11 review · 2026-07 official product page · 2026-07-14 follow-up · 2026-07-14 follow-up · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

Even G2 trades cameras and speakers for productivity and social acceptability

Product: Even Realities G2. Even Realities G2: Even G2 is a camera-free smart-glasses product with a green monochrome heads-up display. Even emphasizes camera-free, lightweight everyday wear; TechCrunch’ s July 11 hands-on notes that it has neither cameras nor speakers and remains heavily dependent on a phone connection. It chooses a different balance among social acceptability, display productivity, and intelligence rather than simply omitting a sensor. This item is handled as confirmed product; missing details remain source not stated.



TechCrunch July 11 hands-on; Even's official G2 page supplies display, weight, IP65, app, and Conversate details.

WHAT IT IS
Even G2 is a camera-free smart-glasses product with a green monochrome heads-up display. Even emphasizes camera-free, lightweight everyday wear; TechCrunch’ s July 11 hands-on notes that it has neither cameras nor speakers and remains heavily dependent on a phone connection. It chooses a different balance among social acceptability, display productivity, and intelligence rather than simply omitting a sensor.

SPECS / STACK
Even lists a 36-gram magnesium-and-titanium frame, 640×350 resolution, 27.5-degree field of view, 60 Hz, 1,200 nits, Micro LED, four microphones, BLE 5.4, IP65, Android and iOS apps, up to two days of battery, and a case with seven full charges. It also lists 35-language translation, 98% passthrough, and prescription support from -12 to +12. TechCrunch reports a \$599 price, no cameras, and no speakers. Chip, model, edge/cloud split, and measured continuous runtime are source not stated.

PAIN POINTS
G2 reduces bystander anxiety associated with camera glasses and moves notifications, captions, and prompts into the line of sight. No speaker avoids open-audio leakage but increases dependence on phone, touch, and display. As the display can carry more content, the product must control long text, notification noise, false activation, and recovery after disconnect; otherwise intelligence becomes an attention burden.

NEW TECH
The increment is a camera-free wearable stack: Micro LED HUD, four microphones, Conversate pre-meeting context and post-meeting summary, 35-language translation, Even Hub, and an optional ring. Even frames no camera, no recording, and lower social friction as design principles and says cloud storage requires explicit consent. R1 combines health tracking with glasses control, creating cross-device input while adding another hardware purchase decision.

LIMITS / UNKNOWNNS
The review exposes unstable phone connectivity, recognition failures in outdoor noise, answers too long to interrupt, manual brightness adjustment in bright rooms, and an R1 price that does not always match its utility. The two-day battery claim depends on display, Conversate, navigation, and notification load; test conditions are not stated. Independent evidence that third-party apps create daily motivation is also missing.

HOW IT WORKS
The user configures the glasses in the Even app, then invokes Even AI, notifications, calendar, QuickList, navigation, translation, Teleprompt, and Conversate through a wake word, temple controls, or Even R1. Even breaks Conversate into Prep Notes, AI Cues, and AI Summary: prepare material before a meeting, receive short context prompts during it, and review a recap afterward. The review observed that long answers stream across the display without a way to skip and that outdoor noise can cause activation and recognition failures.

USE CASES
Supported jobs include loading figures and references into Prep Notes, receiving a definition when a term appears, generating a post-conversation summary, translating speech, reading a teleprompter, checking notifications and calendar, navigating without a phone, and adding tasks. The product fits users who need eye contact, presenters, travelers, and people who do not want to record nearby. The review also says that without meetings, translation, or teleprompting, many users may lack a daily reason to wear it.

USER VOICE
Source not stated.

AVAILABILITY
Even’ s G2 page, support surface, and Even Hub are public. TechCrunch reports a \$599 price and current availability; the R1 is listed at \$249 in the review. Prescription and app information is available, but complete countries, delivery, support, third-party scope, subscription, and enterprise offering are source not stated.

PRODUCT READ
Even G2 makes a coherent bet: remove cameras and speakers to create a text interface that is easier to accept socially. Its hardware and privacy boundary are restrained, but phone connectivity, reading rhythm, and first-party software depth limit usage frequency. It is a thoughtful productivity-glasses product, not yet a universal AI entry point people will wear every day.

COMMUNITY

Community Friction

REVIEW/COMMUNITY FRICTION · REVIEW/COMMUNITY FRICTION · SINGLE-USER ACCOUNT · DOWNGRADED

Community scan: glasses agents are now asked to complete tasks, raising the trust burden

Community

review/community friction · review/community friction · single-user account · downgraded

2026-07-19 Reddit agent-use post · 2026-07-15 WAIC post · 2026-07-21 follow-up

Community scan: glasses agents are now asked to complete tasks, raising the trust burden

Product: AI glasses agent community scan. This scan checked a July 19 Reddit r/augmentedreality post in which a user says Hermes, running through Meta Ray-Bans, created a workout app in a gym and put navigation and a timer into the glasses workflow. This is one user account, not an official capability claim or independent reproduction. The WAIC demo post remains as a contrast for hardware, battery, comfort, and developer-tool gaps.



Reddit community page used as friction evidence: a glasses agent is asked to create an app, navigate, and run a timer; it is one user account, not an official or broadly available capability.

WHAT IT IS
A community-friction scan, not a confirmed product.

SPECS / STACK
The community material provides no Hermes version, Meta API, phone-side execution chain, permissions, chip, OS, price, or shipping evidence, so those implementation details cannot be filled in.

PAIN POINTS
It makes the new friction of moving from asking to executing explicit while retaining the hardware and developer-tool gap list.

NEW TECH
No new technology is confirmed; retain a weak/community signal.

LIMITS / UNKNOWNNS
Single-user account, no independent reproduction, no API or tool-call logs; the next useful evidence is a reproducible demo with permission and undo surfaces.

HOW IT WORKS
One Reddit account says an agent in a Meta Ray-Bans context created a workout app, navigation, and a timer; a separate WAIC post lists hardware, software, AI, battery, comfort, and developer tools as open validation gaps.

USE CASES
It defines validation questions for agent actions on glasses: authorization, state feedback, undo, and failure recovery.

USER VOICE
Source not stated.

AVAILABILITY
The posts describe a personal use account and an exhibition preview, not a public API, retail capability, or reproducible experiment.

PRODUCT READ
Keep it as review/community friction, not product fact.

WILD

Wild Desk

STARTUP SIGNAL · STARTUP SIGNAL · PRODUCT HUNT LISTING
· RESERVATION-ONLY · DOWNGRADED

Monako Glass turns the coding agent monitor into a wearable HUD

CROWDFUNDING SIGNAL · CROWDFUNDING SIGNAL · HANDS-ON REVIEW · OFFICIAL PRODUCT PAGE

MemoMind One trades the camera for a more socially acceptable display-glasses posture

STARTUP SIGNAL · STARTUP SIGNAL · DEVELOPER SURFACE · SOURCE-BACKED PRODUCT PAGE

Brilliant Labs Halo puts open hardware, an edge NPU, and a memory agent behind a developer entry point

Wild

startup signal · startup signal · Product Hunt listing · reservation-only · downgraded

2026-06 Product Hunt launch · 2026-07 to 2026-08 announced shipping window

Monako Glass turns the coding agent monitor into a wearable HUD

Product: Monako Glass. Monako Glass: Monako Glass is a wearable Linux computer and heads-up display for developers. Product Hunt describes a 48-gram device with a waveguide display, bone-conduction microphone, and gesture input running Buildroot Linux. It is not trying to make notifications, music, and photos smaller. It is trying to put the state of Claude Code, Codex, or another coding agent into the wearer’ s field of view after leaving the computer. It remains a reservation-only startup signal, with no independent shipping review on the cited page. This item is handled as startup signal; missing details remain source not stated.



Product Hunt listing: 48g, Buildroot Linux, waveguide HUD, MonoOS, gesture input, \$399 reservation, and shipping window.

WHAT IT IS
Monako Glass is a wearable Linux computer and heads-up display for developers. Product Hunt describes a 48-gram device with a waveguide display, bone-conduction microphone, and gesture input running Buildroot Linux. It is not trying to make notifications, music, and photos smaller. It is trying to put the state of Claude Code, Codex, or another coding agent into the wearer’ s field of view after leaving the computer. It remains a reservation-only startup signal, with no independent shipping review on the cited page.

SPECS / STACK
The public listing supports a 48-gram frame, waveguide heads-up display, bone-conduction microphone, gesture input, a 0.5 TOPS NPU, Buildroot Linux, MonoOS, a Lua app layer, an embedded Rive animation runtime, a 300mAh battery, approximately four hours of screen-on use, approximately eight hours of normal use, and support for Claude Code, Codex, and other coding agents. Product Hunt also lists a \$19 reservation toward a \$399 unit and a July-to-August 2026 shipping window. Display resolution, field of view, camera, RAM, storage, CPU, agent API, permission sandbox, edge/cloud model, and security-update policy are source not stated.

PAIN POINTS
Monako Glass targets a specific coding-agent problem: an agent may run for a long time, but the developer must keep watching a laptop to know whether it is progressing, failing, or waiting for confirmation. A HUD moves waiting and state feedback from a fixed monitor closer to the body, while the bone-conduction microphone aims to preserve environmental hearing. The cost is putting a dense engineering workflow into a narrow and potentially distracting view. If output is too long, state is opaque, or errors cannot be undone, visibility becomes anxiety rather than control.

NEW TECH
The distinctive technical story is the combination of a Linux device, waveguide HUD, 0.5 TOPS NPU, gesture input, and a Lua application runtime. Letting an agent generate and run a Lua app could make the application layer a dynamic artifact and reduce pre-build and installation steps; the Rive runtime supplies an animation basis for status feedback. This moves coding agents from IDE plugins toward a wearable system surface. The public material does not prove that a device-side agent can safely generate, sign, and run arbitrary apps, nor does it show permission, network, code-secret, or update isolation.

LIMITS / UNKNOWNNS
The key validation questions are whether 48 grams includes the full frame and battery; how four hours of screen-on and eight hours of normal use were measured; what the 0.5 TOPS NPU actually does; whether Lua apps really run reliably without a build step; whether Claude Code and Codex connect locally, remotely, or through the user’ s existing computer; whether the HUD can show long logs, diffs, and security confirmations; and how the system pauses, undoes, and recovers from agent failure. Product Hunt comments raise source-licensing, latency, and feasibility questions. Those are community friction signals, not official failure findings.

HOW IT WORKS
The intended workflow is to describe a task through voice or the bone-conduction microphone, use gesture or hand tracking to inspect agent state, and read short output, waiting, error, or confirmation states in the waveguide HUD. Product Hunt says MonoOS provides a Lua application layer where an agent can generate and run a Lua app immediately without a traditional build step. The conceptual loop is describe a task, let the agent generate the surface, run it in view, and continue or confirm. Wake phrase, gesture vocabulary, code-execution permissions, network connection, rollback, and companion-computer behavior are source not stated.

USE CASES
The immediate job is for a developer who wants to leave the desk while still knowing whether an agent completed, stalled, needs input, or produced an error. In a lab, exhibition hall, or moving workday, the HUD could monitor a long task and allow a voice or gesture follow-up. The Lua app concept could let an agent generate a lightweight task-specific interface instead of sending the user back to an IDE and another window. This is an agent status panel and control surface, not a complete code editor. Long diffs, permission review, secret entry, and final code inspection still need a larger screen.

USER VOICE
Source not stated.

AVAILABILITY
Product Hunt lists a \$19 reservation that can be applied toward a \$399 unit and a July-to-August 2026 shipping window; it also marks the product reservation-only with no reviews. The official Monako site is a JavaScript application with limited independently visible technical specifications and support policy. There is no verifiable evidence in the cited sources for production inventory, regional support, returns, an SDK download, a public source repository, an OS version, or real user shipping reviews. It should not be described as broadly available or as a validated developer tool.

PRODUCT READ
Monako Glass is a clear startup signal. It does not repackage the claim that AI glasses can take photos; it builds around the waiting and feedback problem of long-running agents. The 48-gram Linux/HUD/Lua/coding-agent combination is a useful product hypothesis, but reservation status, no independent review, and missing security surfaces require a downgrade. Product teams can study how agent state might be understood while moving, but should not treat this as a validated production coding surface.


Wild

crowdfunding signal · crowdfunding signal · hands-on review · official product page

2026-06-28 Kickstarter · 2026-07 Android Central hands-on

MemoMind One trades the camera for a more socially acceptable display-glasses posture

Product: XGIMI MemoMind One. XGIMI MemoMind One: MemoMind One is a camera-free display AI glasses product from XGIMI’ s MemoMind brand. Its core job is not to see the world, but to put private display, open-ear audio, and AI widgets into a frame that looks closer to everyday eyewear. The official page lists about 46.6 grams, dual Micro-LED projection, 2,000 nits, an adjustable one-to-five-metre virtual distance, reduced light leakage, Harman AudioEFX, and triple directional microphones. Reviews place it in concrete flows for notifications, calendar, to-do lists, voice notes, translation, and walking or cycling navigation. This item is handled as crowdfunding signal; missing details remain source not stated.



Android Central hands-on page: camera-free design, 46g weight, notifications/calendar/to-do/translation, and Kickstarter delivery timing.

WHAT IT IS
MemoMind One is a camera-free display AI glasses product from XGIMI’ s MemoMind brand. Its core job is not to see the world, but to put private display, open-ear audio, and AI widgets into a frame that looks closer to everyday eyewear. The official page lists about 46.6 grams, dual Micro-LED projection, 2,000 nits, an adjustable one-to-five-metre virtual distance, reduced light leakage, Harman AudioEFX, and triple directional microphones. Reviews place it in concrete flows for notifications, calendar, to-do lists, voice notes, translation, and walking or cycling navigation.

SPECS / STACK
The official product page lists three frame styles, beta-titanium, magnesium-aluminium, or acetate materials, 46.6 grams, dual Micro-LED projectors, 2,000 nits, a one-to-five-metre adjustable display distance, downward light and a high-transparency grating, no camera, Harman AudioEFX open-ear audio, triple directional microphones, and ZEISS XRRX prescription lens options. Android Central confirms dual Micro-LED waveguides, built-in audio, and 46 grams. Battery capacity, processor, OS/API, display resolution, edge/cloud model split, and measured runtime are source not stated.

PAIN POINTS
MemoMind One targets two frictions. Camera glasses are difficult to trust in public, while traditional HUD products often require a phone or dedicated controller. Removing the camera makes the privacy boundary easier to explain; dual Micro-LED puts short prompts back into the visual field; voice notes reduce the effort of saving an idea. The trade-off is a narrower context window: the glasses cannot capture visual context, navigation starts with phone setup, and the AI experience depends on configuration and connectivity.

NEW TECH
The product combination is camera-free hardware, dual Micro-LED waveguides, 2,000-nit display, adjustable one-to-five-metre virtual distance, low-leakage optics, Harman AudioEFX, and a widget-oriented AI experience. The innovation is not a stronger vision model. It is treating ‘I do not record you’ as a form-factor decision and using display to make notifications, translation, and navigation a glanceable layer. Social acceptance, display privacy, and all-day wearability become one product chain.

LIMITS / UNKNOWNNS
The main unknown is not whether it can display. It is whether the delivered product remains useful over time. Battery life under mixed display, audio, microphone, translation, and notification load is unverified; phone-disconnected behavior is unclear; leakage at different viewing angles needs testing; and the storage and deletion model for voice notes is not fully disclosed. Reviews also note that navigation is limited to walking and cycling, while AI response time and audio quality require continued observation.

HOW IT WORKS
The user configures the glasses through the MemoMind app, mirrors notifications, calendar, news, to-do items, and an idea board, dictates notes, and invokes audio recording or real-time translation. Navigation requires entering a destination in the app first, then presenting walking or cycling turn-by-turn cues. Because there is no camera, the product cannot directly see a menu, sign, or gesture; visual understanding and capture are outside its loop. It trades active configuration, short display cues, and whispered audio for a clearer camera boundary.

USE CASES
The device is for glancing at notifications, calendar, news, and to-dos; dictating an idea while moving; listening to audio; translating speech in real time; and receiving navigation while walking or cycling. A camera-free design reduces bystander concern about invisible recording and makes the product easier to read as normal eyewear. For users who need visual AI, it is not a replacement for camera glasses. The product choice is explicit: it prioritizes prompts and memory capture over seeing and interpreting the world in front of the wearer.

USER VOICE
Source not stated.

AVAILABILITY
Android Central reports a Kickstarter campaign with a \$399 early-bird pledge, a \$499 prescription option, customized designs from \$449, and expected delivery in August 2026. Tom’ s Guide reports a possible \$599 commercial price. The official page exposes a reservation path and specifications, but a crowdfunding promise is not current retail inventory. Regions, batch delivery, warranty, subscription, API access, and final retail pricing remain source not stated and should be rechecked after the campaign.

PRODUCT READ
MemoMind One turns the social problem of AI glasses into a product strategy: remove the camera, keep display, audio, and configurable widgets. It does not try to see everything, which narrows the job but may improve public acceptance. Keep the label at crowdfunding signal: the hardware direction is concrete and early hands-on evidence exists, while delivery, battery, phone dependence, and long-term software support have not yet been validated at consumer scale.

startup signal · startup signal · developer surface · source-backed product page

Brilliant Labs Halo puts open hardware, an edge NPU, and a memory agent behind a developer entry point




Brilliant Labs

Contact

News

Dev

ORDER NOW



Halo

\$349.00 USD

Introducing Halo!

Open source AI glasses for the curious, creative, and forward thinking.

Halo features a fresh design, reimaged optics and electronics, and Noa, your private conversational AI agent with long-term memory of your life. With Minispeech, Halo lets you build new experiences using natural language and share them with others on our App Store.

These features are included free out of the box with daily usage caps.

-all while weighing just over 40g

Shipping starts the week of July 20.
Secure Halo at \$349 before the price increases to \$399.

COLOR

Black

QUANTITY

- 1 +

BUY NOW

Import taxes may apply at delivery and vary by country.

SOURCES [Brilliant Labs Halo product page](#) [Brilliant Labs developer page](#) [Halo developer documentation](#) [Brilliant Labs and Alif partnership](#) [Brilliant community shipping discussion](#)

Halo is Brilliant Labs' open AI glasses for creators and developers. The official page lists a \$349 price and combines Noa, a private conversational agent with long-term memory, Vibe Mode, a color display, open hardware, and open software. The bet is that glasses become a programmable embodied interface rather than a closed assistant that users merely consume.

Brilliant lists a little over 40 grams, a Micro Color OLED, two bone-conduction speakers, an Alif Balletto B1 low-power AI processor with a Cortex-M55 CPU and Ethos-U55 NPU, a low-power optical sensor, dual microphones, a six-axis IMU, tap detection, Bluetooth 5.3, ZephyrOS with a Lua interface, a cross-platform mobile app, a cloud AI agent, and a 14-hour all-day battery claim on its specifications page. The Alif partnership supports the edge-AI direction; RAM, camera details, model parameters, and edge/cloud split are source not stated.

An open platform lowers the barrier created by closed glasses that cannot be modified, connected to a preferred model, or used for sensor research. The edge-NPU route tries to reduce dependence on the cloud for continuous visual and audio tasks. Long-term memory reduces the need to restate context. The cost is that users take on privacy, model choice, app quality, and hardware maintenance, while memory agents create stronger requirements to explain and delete what the device remembers.

There is an evidence gap between the published specification and a mass-produced experience. Independent testing has not established that 14 hours survives mixed display, microphone, memory, and cloud-request load. Public sources do not state the NPU models, camera capability, false-wake rate, measured data retention, or a developer marketplace. Brilliant's terms say it does not currently operate a public developer marketplace, and community posts report shipping uncertainty.

A user can wake Noa by voice or touch and talk about what they see, hear, or imagine; Vibe Mode lets a person describe an experience in natural language. Developers use the Brilliant SDK, Flutter SDK, and documentation to build applications, while open hardware and software repositories provide deeper modification paths. The product page describes memory enhancement, visual and audio dialogue, and translation, while exact gestures, display feedback, third-party agent permissions, and distribution remain to be confirmed in developer material.

Halo is positioned for connecting objects seen, speech heard, and personal memory, as well as translation, hands-free questions, small HUD applications, creative prototypes, and experimental agents. A developer can combine the optical sensor, microphones, IMU, display, and bone-conduction audio into a custom feedback loop. For users, the promise is remembering what they encountered; for creators, the value is a shared experimental surface for hardware, SDKs, Lua/Flutter, and community projects.

Source not stated.

Halo's product page lists \$349 and says shipping starts in Q1 2026. The developer page, documentation, and product page are public; prescription and sunglass options are offered through a partner, with the display optic listed as adjustable from +2 to -6 diopters. The current page does not state inventory and delivery for every country, Noa+ coverage, a complete app store, or general SDK eligibility. Community discussion still questions actual shipping and wait times, so this remains a startup signal.

Halo represents a developer-first AI-glasses route that hands the energy budget, source access, sensors, and agent memory to experimenters. The \$349 price lowers the cost of trying the form factor, while open design and an edge NPU make it valuable for research. Production, shipping, distribution, and memory governance decide whether it is a usable platform or an attractive kit. Keep the verdict at startup signal rather than mature consumer ecosystem.

RESEARCH

Research Watch

RESEARCH SIGNAL · RESEARCH SIGNAL · OPEN HARDWARE
PROTOTYPE · DOWNGRADED

Research scan: OpenGlass makes event-based vision and power management a reusable prototype

Research research signal · research signal · open hardware prototype · downgraded

2026-06-05 arXiv submission · 2026-06-08 revision

Research scan: OpenGlass makes event-based vision and power management a reusable prototype

Product: OpenGlass research prototype. OpenGlass is an arXiv research signal submitted June 5 and revised June 8, 2026, not a shipped glasses product. The paper exposes a modular FPC interposer, nRF5340 event-driven wake-up, a GAP9 RISC-V SoC, a Prophesee GENX320 event camera, a 200mAh battery, and open hardware, firmware, and models.

arXiv

Computer Science > Computer Vision and Pattern Recognition

Submitted on 5 Jun 2026 (v1), last revised 8 Jun 2026 (this version, v2)

OpenGlass: Ultra-Low-Power On-Device AI Eyewear with Event-based Vision

Pietro Bonazzi, Julian Moosmann, Ahmet Celik, Philipp Mayer, Michele Magno

Smart eyewear enables unobtrusive, context-aware interaction through multimodal sensors and on-device intelligence, but is severely limited by power, memory, and compute constraints in a compact form factor. Open-hardware platforms supporting event-based vision and embedded ML at this scale are rare. This work introduces an open-source smart-glasses platform for rapid prototyping of novel sensors and algorithms. Its modular design uses a flexible FPC interposer to support both event-based and frame-based cameras without full PCB redesign. A hardware-software co-designed power management system combines a configurable PMIC with event-driven wake-up via an nRF5340 coordinator, keeping the GAP9 RISC-V SoC powered down between inferences. The prototype achieves up to 11.5 hours of continuous on-device ML from a 200 mAh battery. As a demonstration, an egocentric hand gesture recognition pipeline was evaluated on the Lynx dataset using polarity-separated event histograms from a Prophesee GENX320 camera. R12+1D achieved the best cross-subject accuracy of 83.94% (macro F1 = 0.781) under leave-two-subjects-out validation, with 78.3 ms end-to-end inference latency on the GAP9. Temporal augmentation and removal of ambiguous classes provided the largest gains (+8.9 pp). All hardware designs, firmware, and models are released open source.

Subjects: Computer Vision and Pattern Recognition (cs.CV)

Cite as: arXiv:2606.07431 [cs.CV] (or arXiv:2606.07431v2 [cs.CV] for this version) (https://doi.org/10.48550/arXiv.2606.07431)

Submission history

From: Pietro Bonazzi [view email]

[v1] Fri, 5 Jun 2026 16:27:02 UTC (17,492 KB)

[v2] Mon, 8 Jun 2026 14:50:28 UTC (17,517 KB)

Access Paper:

[View PDF](#)

[HTML \(experimental\)](#)

[TeX Source](#)

[\[PDF\] \[HTML\] \[TeX\]](#)

Current browse context: cs.CV

[< prev](#) | [next >](#)

[new](#) | [recent](#) | [2026-06](#)

Change to browse by: cs

References & Citations

[NASA ADS](#)

[Google Scholar](#)

[Semantic Scholar](#)

[Export BibTeX Citation](#)

Bookmark

Bibliographic Tools

Code, Data, Media

Demos

Related Papers

About arXivLabs

Bibliographic and Citation Tools

☐ Bibliographic Explorer (What is the Explorer?)

☐ Connected Papers (What is Connected Papers?)

☐ Litmaps (What is Litmaps?)

☐ scite Smart Citations (What are Smart Citations?)

Which authors of this paper are endorsers?

Disable MathJax (What is MathJax?)

arXiv paper page; event-based vision, power management, and open-source scope are research evidence.

WHAT IT IS

An open smart-glasses research platform, not a consumer product.

SPECS / STACK

FPC interposer, nRF5340, GAP9 RISC-V, Prophesee GENX320, 200mAh battery, and the LynX dataset; the paper reports up to 11.5 hours of continuous edge ML, 83.94% accuracy, and 78.3ms latency.

PAIN POINTS

It separates the power and compute bottleneck of continuous sensing into event wake-up and edge inference problems.

NEW TECH

Event-based vision, a modular camera interface, and hardware-software co-designed power management.

LIMITS / UNKNOWNNS

Mass production, all-day wear, privacy, false wakes, and real-world performance are unverified; keep every claim at research-signal level.

HOW IT WORKS

An event camera detects change, the nRF5340 wakes the GAP9 for on-device gesture inference, and the system sleeps between inferences; the paper does not provide a complete daily-wear interaction flow.

USE CASES

Gesture-recognition research, open sensor experiments, and low-power visual interaction.

USER VOICE

Source not stated.

AVAILABILITY

Paper, hardware designs, firmware, and models are public; there is no consumer shipping, price, warranty, or support evidence.

PRODUCT READ

Keep as a research watch item, never as a substitute for shipping-product evidence.

PATENT

Patent Watch

PATENT SIGNAL · PATENT SIGNAL · EXPLICITLY SPECULATIVE


Patent lane: smart-glasses IP is active, but it signals direction rather than product evidence

Patent patent signal · patent signal · explicitly speculative

2026 patent/source scan · 2026-07-11 follow-up · 2026-07-12 follow-up · 2026-07-13 follow-up · 2026-07-14 follow-up · 2026-07-14 follow-up · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

Patent lane: smart-glasses IP is active, but it signals direction rather than product evidence

Product: Smart-glasses patent lane scan. Smart-glasses patent lane scan: A patent-signal scan. WIPO's Rokid material shows how smart-glasses companies use patents, trademarks, shipments, and deployments to build global credibility. Google Patents documents around AI-supported smart glasses and a GS1 design patent show possible medical and form-factor directions. Android Central's coverage of the Solos lawsuit against Meta and EssilorLuxottica shows that IP friction may affect the smart-glasses market. None of this is treated as a shipped feature.. This item is handled as patent signal; specs, price, regions and dates use cited source claims only, and missing details remain source not stated.



Source-backed screenshot from Google Patents, explicitly treated as patent signal rather than product evidence.

WHAT IT IS
A patent-signal scan. WIPO's Rokid material shows how smart-glasses companies use patents, trademarks, shipments, and deployments to build global credibility. Google Patents documents around AI-supported smart glasses and a GS1 design patent show possible medical and form-factor directions. Android Central's coverage of the Solos lawsuit against Meta and EssilorLuxottica shows that IP friction may affect the smart-glasses market. None of this is treated as a shipped feature.

SPECS / STACK
The sources support WIPO's statements about Rokid's weight, patent count, trademarks, countries, and deployment examples; Google Patents supports document titles, intended use, abstracts, and publication records; Android Central supports the reporting context for Solos suing Meta and EssilorLuxottica. Claim validity, court outcomes, actual infringement, product injunction risk, licensing fees, whether a patent is implemented, and final device specs are source not stated.

PAIN POINTS
Patent scanning addresses a blind spot in product research. Teams often track announcements and reviews but miss the way a hardware category is shaped by patents and litigation before launch. Smart glasses are especially exposed because optics, audio, sensing, exterior design, and privacy indicators may all be constrained. Reading IP signals early prevents teams from treating a legally or technically constrained component as a low-risk design foundation.

NEW TECH
The signal is not one confirmed feature. It is the clustering of protected directions: lighter glasses, medical and industrial assistance, AI-supported diagnosis, exterior design, audio/sensing architecture, and display paths. Smart-glasses competition is expanding from device specs into IP portfolios and ecosystem licensing.

LIMITS / UNKNOWNNS
Patent language is broad, translations can be rough, claims are complex, and implementation status is often unknown. Many patents are defensive and never appear in products. Lawsuit reporting also changes as courts move. This issue treats the lane as watchlist only, not confirmed product evidence.

HOW IT WORKS
A patent has no user workflow. It can only indicate components and directions future products might emphasize: medical assistance, industrial use, form factor, sensing, audio, display, waveguides, environmental understanding, or multimodal input. Lawsuit coverage suggests possible indirect user impact such as sales risk, brand partnership changes, licensing costs, or feature constraints. These are external signals and cannot replace official product pages or hands-on testing.

USE CASES
The role of the patent lane in a product magazine is directional and risk-oriented. It helps identify which interaction capabilities are being fenced, which hardware forms might become competitive barriers, and which lawsuits could influence channel, price, or availability. For product teams, it is a reminder that launch coverage is not enough. IP, standards, supply chain, and legal risk can shape the roadmap before users see a device.

USER VOICE
Patent documents do not contain user voice. The user impact of lawsuits is also speculative at this stage. The cited sources do not show consumers changing purchase behavior directly because of these patent documents; source not stated.

AVAILABILITY
Patent publication does not equal product launch. A filed lawsuit does not equal a court decision. WIPO profile material does not mean the same product is purchasable in every region. Availability must be verified through specific brand stores, official developer docs, or regulatory notices.

PRODUCT READ
The patent-lane read is deliberately conservative: smart-glasses IP is dense enough that product managers should include patents and litigation in competitive scanning, but user value still has to be proven through purchasable products, real interactions, and source-backed claims. No patent in today's issue is presented as a launch fact.

CHINA

China / Global Compare

CONFIRMED PRODUCT · CONFIRMED PRODUCT · OFFICIAL
PRODUCT PAGE · CHINA PRODUCT REPORT

MOONIX makes the AI glasses case start with
normal wearability

DEVELOPER SURFACE · DEVELOPER SURFACE · CHINA/GLOBAL
PRODUCT SIGNAL

Z.ai ZCode packages GLM-5.2 as a China/global
developer-tool signal

CONFIRMED PRODUCT · CONFIRMED PRODUCT · CHINA
LAUNCH COVERAGE · COMPANY ANNUAL-REPORT
DISCLOSURE

iFlytek AI Glasses turn translation, prompting,
and GlassClaw into a China-market workflow

CONFIRMED PRODUCT · CONFIRMED PRODUCT · CHINA
PRODUCT SIGNAL · OVERSEAS RETAIL EVIDENCE

Rokid uses YodaOS to push AI glasses from
phone accessory toward an AIOS entry point

China

confirmed product · confirmed product · official product page · China product report

2026-07 official MOONIX product page · 2026-07-17 IT之家 launch report

MOONIX makes the AI glasses case start with normal wearability

Product: MOONIX AI Glasses. MOONIX AI Glasses: MOONIX is an everyday audio and AI glasses product from Xinmou Technology. The standard model is presented around approximately 14.9 grams, about 16 hours of all-day battery, proactive AI, audio, microphones, translation, reminders, a voice assistant, and AI capture plus summary. It does not try to put display, navigation, and visual understanding into the first product surface. It starts with whether people can wear the device every day and whether important information can be retained. The official FAQ distinguishes the standard and Pro models: Pro adds a high-definition camera and audio-video recording, while the standard model focuses on audio capture and AI questions. This item is handled as confirmed product; missing details remain source not stated.



MOONIX official page: 14.9g, 16h, AI capture and summary, six-microphone array, and everyday scenarios.

WHAT IT IS

MOONIX is an everyday audio and AI glasses product from Xinmou Technology. The standard model is presented around approximately 14.9 grams, about 16 hours of all-day battery, proactive AI, audio, microphones, translation, reminders, a voice assistant, and AI capture plus summary. It does not try to put display, navigation, and visual understanding into the first product surface. It starts with whether people can wear the device every day and whether important information can be retained. The official FAQ distinguishes the standard and Pro models: Pro adds a high-definition camera and audio-video recording, while the standard model focuses on audio capture and AI questions.

SPECS / STACK

The official product page supports approximately 14.9 grams for the standard model, about 19.9 grams for Pro, an advertised 16-hour all-day battery, a six-microphone array, audio, proactive AI, iPhone/iPad and Android apps, and a firmware-update entry point. The FAQ states that the standard model emphasizes audio capture and AI questions while Pro adds a high-definition camera and audio-video recording. The processor, RAM, storage, OS/API version, network protocol, camera model, audio codec, cloud model, edge-inference split, and charging structure are not published on the cited pages. IT之家 reports global sale from July 17, 2026, with official pricing starting at RMB 2,299; configuration and inventory remain region dependent.

PAIN POINTS

MOONIX targets a problem beyond the lack of AI on a phone: a phone must be pulled out, unlocked, and navigated to a recording entry before the moment is saved. At 14.9 grams, the product directly targets heavy, thick-templed, visibly technological smart glasses. Audio and microphones move capture to the face, AI summaries reduce post-meeting cleanup, and the 16-hour claim makes all-day wear the intended frame. The standard-versus-Pro split also makes privacy and capability a product choice. The tradeoff is equally clear: without a display or camera, the standard model cannot put information back into the field of view or understand a menu, sign, or object in front of the user.

NEW TECH

The public evidence supports a combination of an ultra-light structure, a six-microphone array, proactive AI, voice capture, AI summary, and a continuous layer connected to a phone app, firmware, and a longer-term personal knowledge base. The innovation is an interaction allocation: camera, display, and complex visual jobs are left to the higher-tier model or another device, while the standard device concentrates its budget on listening, capture, reminders, and review. ‘The more you use it, the more it understands you’ depends on persistent information accumulation. The key unanswered questions are when recordings upload, how they are classified, how they are deleted, and whether they can be exported. Data lifecycle and API details are source not stated.

LIMITS / UNKNOWNNS

The next validation target is not the lightweight number alone but the daily loop: six-microphone false capture in trains, group meetings, and wind; summary latency, misremembered content, and speaker attribution; whether 16 hours holds under AI, Bluetooth, continuous capture, and high-volume mixed load; how the standard model shows recording, pause, and deletion; what remains when the phone or cloud service is unavailable; and how the personal knowledge base is searched, exported, and cleared. The official FAQ points users to privacy policy and agreements for camera, microphone, and cloud AI handling. There is not enough evidence to infer default retention.

HOW IT WORKS

The user first pairs the glasses through the iPhone, iPad, or Android companion app and manages AI capabilities and firmware updates there. During wear, voice and the glasses’ control entry points start capture, questions, translation, reminders, and review. The official page frames a daily rhythm of 9:00 inspiration capture, an 11:00 lecture, a 15:00 business conversation, and a 20:00 information review. Short recordings are turned into AI summaries and a longer-term personal knowledge base. The standard path is hear, capture, summarize, and review rather than see, overlay, and navigate. Exact buttons, wake phrase, summary latency, and failure recovery are source not stated.

USE CASES

Concrete jobs are inspiration capture, meetings, lectures, interviews, and commuting review. The wearer can save a spoken thought before taking out a phone, then use AI to extract the important points and build a persistent personal knowledge base. Translation, reminders, and voice assistance cover shorter interventions. The six-microphone array is positioned as a way to capture speech in noisy environments, making ‘save first, organize later’ the central workflow. The camera-free standard model fits people who prioritize appearance, low weight, and low interruption. A person who needs photos or visual understanding must choose Pro or another camera-equipped product.

USER VOICE

Source not stated.

AVAILABILITY

MOONIX’ s official site presents standard and Pro products, downloads, help, privacy, and user-agreement entries. IT之家 reports that the first product went on global sale on July 17, 2026, with official pricing from RMB 2,299. The site exposes iPhone/iPad and Android download routes, but the cited pages do not provide a complete list of country inventory, taxes, support, subscriptions, regional AI-service availability, prescription options, or delivery windows. Exact Pro and standard pricing, frame combinations, camera model, runtime differences, and sales channels should be checked at checkout and in the user guide; missing details remain source not stated.

PRODUCT READ

MOONIX makes a coherent product bet: establish weight, appearance, capture, and review as a daily habit before adding more visible intelligence. Its entry point differs from display and camera glasses. The 14.9-gram claim and camera boundary make social acceptance easier to discuss, while narrowing the capability surface. As a confirmed product, it is worth watching for whether people repeatedly say ‘remember this’ and trust that the recording is visible, deletable, and retrievable. Long-term data governance, accuracy, and disconnected behavior still require hands-on validation.

China

confirmed product · confirmed product · China launch coverage · company annual-report disclosure

2026-05-28 launch · 2026-05-29 36Kr product report

iFlytek AI Glasses turn translation, prompting, and GlassClaw into a China-market workflow

Product: iFlytek AI Glasses + GlassClaw. iFlytek AI Glasses + GlassClaw: iFlytek AI Glasses are display-enabled multimodal AI glasses launched by iFlytek in Macau on May 28, 2026, positioned as a ‘super AI assistant in front of your eyes.’ The 36Kr launch report lists an approximately 40-gram frame, full-scenario translation, multimodal noise reduction, prompting, and GlassClaw. iFlytek’s annual report adds self-developed simultaneous interpretation and multilingual models, AI visual and speech translation, lip-motion noise reduction, and a May 28 market launch. This item is handled as confirmed product; missing details remain source not stated.



36Kr launch report: 40g body, translation, five air-conduction mics plus bone conduction and lip reading, GlassClaw, price, and retail channels.

WHAT IT IS
iFlytek AI Glasses are display-enabled multimodal AI glasses launched by iFlytek in Macau on May 28, 2026, positioned as a ‘super AI assistant in front of your eyes.’ The 36Kr launch report lists an approximately 40-gram frame, full-scenario translation, multimodal noise reduction, prompting, and GlassClaw. iFlytek’s annual report adds self-developed simultaneous interpretation and multilingual models, AI visual and speech translation, lip-motion noise reduction, and a May 28 market launch.

SPECS / STACK
Public sources support an approximately 40-gram magnesium-aluminium frame, resin diffractive waveguide lenses, a custom micro-optical engine, five air-conduction microphones plus one bone-conduction microphone, lip-motion recognition, multimodal noise reduction, binocular monochrome display, GlassClaw, AstronClaw architecture, and iFlytek multilingual models. The 36Kr report says real-time translation across 122 languages. The standard model is priced at RMB 4,299 and the longer-battery model at RMB 4,699. Processor, RAM, display resolution, camera model, runtime in hours, recording retention, and edge/cloud split are source not stated.

PAIN POINTS
It targets multiple-speaker overlap, tool switching during translation, separate prompting and note-taking systems, and the manual cleanup required after a meeting. Lip-motion recognition and sound-source selection turn ‘who is speaking?’ into a system input; binocular display puts translation and prompts into view; GlassClaw extends capture into a proposal and email. The risk is direct: selecting the wrong speaker, delayed translation, or an overreaching agent can inject errors into a business interaction.

NEW TECH
The product combines multimodal noise reduction with an agent workflow. Microphones, bone conduction, lip motion, and vision help select a target, after which interpretation, visual translation, or GlassClaw acts. GlassClaw is framed as more than Q&A: it supports capture, retrieval, document generation, and email delivery. The HCI problem therefore becomes continuous: who is selected, what is heard, what is shown, what is generated, and when it is sent.

LIMITS / UNKNOWNNS
The evidence is mainly a launch event, media demonstrations, and company reporting; long-term independent review is missing. Unknowns include language-by-language quality across 122 languages, wrong-speaker rate, lip recognition with masks, profile faces, or low light, outdoor display legibility, mixed-load runtime during translation plus GlassClaw, confirmation and rollback for generated content, and permission UI before cross-device sending. A demonstration proves that the flow exists, not that it is reliable in a real business setting.

HOW IT WORKS
A wearer uses voice, gaze, and device controls during a meeting, exhibition, or call to start simultaneous interpretation, face-to-face translation, online interpretation, call translation, visual translation, prompting, meeting notes, and task processing. The reported interaction is ‘who to look at, who to listen to, who to translate’: five air-conduction microphones, one bone-conduction microphone, sound-source localization, visual recognition, and lip-motion recognition help select the target speaker. GlassClaw can use what the glasses have seen and heard to retrieve information, generate a proposal or email, and hand work across devices in the demonstrated flow.

USE CASES
The product focuses on cross-language business: face-to-face translation at an exhibition or airport, prompts for meetings and presentations, call translation, interview notes, visual translation, target-speaker capture in noisy rooms, and extracting information from an event poster to generate a partnership proposal and email. It connects ‘hear, translate, record, generate, distribute’ into one workflow for people who communicate across languages and cannot keep looking down at a phone.

USER VOICE
Source not stated.

AVAILABILITY
36Kr reports a RMB 4,299 standard model and RMB 4,699 longer-battery model, reservations from May 28, full-price 618 presale from June 15, and about 1,000 Chinese stores offering trial, purchase, and prescription fitting. iFlytek’s annual report says the product was scheduled to launch in Macau on May 28, 2026. International sales, inventory, the exact battery-model difference, subscriptions, API access, retention, and enterprise administration remain source not stated.

PRODUCT READ
iFlytek AI Glasses are one of the most complete China-market workflow products in this scan: translation, prompting, notes, and the GlassClaw agent sit in one roughly 40-gram display frame, with public price and retail channels. The real test is not feature count. It is whether the continuous state of who is heard, what is translated, what is generated, and who receives it remains reviewable and correctable. Treat it as a confirmed product with experience quality and edge/cloud details still to validate.

China

confirmed product · confirmed product · China product signal · overseas retail evidence

2026-06-29 YodaOS announcement · 2026-07-10 Japan general sale · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

Rokid uses YodaOS to push AI glasses from phone accessory toward an AIOS entry point

Product: Rokid Smart AI Glasses + YodaOS. Rokid Smart AI Glasses + YodaOS: Rokid Smart AI Glasses are AI/AR glasses with cameras, microphones, speakers, and binocular display. YodaOS is the AI-native smart-glasses operating system Rokid introduced publicly in June 2026. The China signal has moved beyond one pair of glasses toward an OS, model switching, third-party services, and an agent store: Rokid says users can choose among ChatGPT, Gemini, DeepSeek, and Alibaba’ s Qwen, positioning the glasses as an AI entry point rather than a phone accessory. This item is handled as confirmed product; missing details remain source not stated.



Rokid official page showing display AI glasses, developer services, and YodaOS-Master; specifications and Japan availability follow product and review sources.

WHAT IT IS

Rokid Smart AI Glasses are AI/AR glasses with cameras, microphones, speakers, and binocular display. YodaOS is the AI-native smart-glasses operating system Rokid introduced publicly in June 2026. The China signal has moved beyond one pair of glasses toward an OS, model switching, third-party services, and an agent store: Rokid says users can choose among ChatGPT, Gemini, DeepSeek, and Alibaba’ s Qwen, positioning the glasses as an AI entry point rather than a phone accessory.

SPECS / STACK

Mogura VR lists approximately 49 grams, Sony IMX681 12MP, a 109-degree wide-angle lens, F2.25, HDR, noise reduction and stabilization, binocular Micro LED displays, 1,500 nits, a 30-degree field of view, Snapdragon AR1 Gen 1, Wi-Fi 6, Bluetooth 5.3, 2GB RAM, and 32GB ROM for the Japan product. The system also includes YodaOS, a phone app, AI services, display, and audio output. Rokid’ s official page shows YodaOS-Master and developer services but does not state every model’ s price, battery, retention policy, or edge/cloud split; missing details remain source not stated.

PAIN POINTS

Rokid targets dependence on a phone and a single model. A user can trigger a capability by intent rather than opening an app; partners can place an agent in a platform rather than rebuild each feature per device; overseas users can choose models and channels by region. Multimodal input targets walking, travel, and work where staring at a phone is awkward. The cost is a harder explanation problem across models, permissions, quality, and accountability.

NEW TECH

YodaOS makes an ‘AI-native OS’ an intent-management and agent-ecosystem proposition, not merely a chatbot in settings. Rokid highlights model switching, integrations with WeChat, Douyin, Alipay, JD.com, Alibaba Cloud, Google Cloud, and AWS, plus an overseas agent store. Display, camera, voice, movement, and AI services become one multimodal entry point. The technical moat depends on routing, context persistence, permission confirmation, offline fallback, and reproducible developer tools.

LIMITS / UNKNOWNNS

YodaOS has a strong public narrative, but the complete SDK, agent permission model, app review, token economics, network-failure behavior, privacy indicators, and third-party data paths are not all public. The detailed specifications come from a Japan interview report and should not be generalized to every Rokid model. Retention, daily use, translation accuracy, outdoor readability, and battery need independent measurement. Calling it the industry’ s first AI-native OS as a universal fact would exceed the evidence.

HOW IT WORKS

The wearer uses voice, temple touch, gaze, or body movement to trigger photos, translation, AR navigation, visual questions, and tasks. The launch material describes a camera capture, speech recognition, model processing, display or audio feedback loop. YodaOS connects multimodal input to agents and external services: translate a foreign sign, provide a travel route, or call a work procedure. The overseas agent platform has been released and an agent store was scheduled for July; exact catalog and rollout need verification.

USE CASES

Reported jobs include face-to-face and travel translation, reading signs and menus, AR navigation, hands-free capture, visual questions, work prompts, and services involving payments, livestreaming, or cloud tools. The Japan product is described as an all-in-one camera, translation, and AR-navigation device and went on general sale on July 10, 2026 through Rokid’ s store, Amazon, Rakuten, and 75 physical stores. The China/global distinction is which models, services, and retail path are locally available.

USER VOICE

Source not stated.

AVAILABILITY

Rokid’ s China site lists display AI glasses, developer services, and YodaOS-Master. The Japan report confirms general sale from July 10, 2026 at 109,990 yen including tax through the official store, Amazon, Rakuten, and 75 physical stores. China pricing by model, the final overseas agent-store catalog, regional model availability, prescription options, delivery, and support are source not stated. Crowdfunding totals are platform/company signals, not long-term active-user evidence.

PRODUCT READ

Rokid is the most complete China/global AIOS signal in today’ s issue: the hardware is in Japanese general sale, while the OS, model switching, and agent-store story tries to move the ecosystem from device sales to a service entry point. Its opportunity is openness and localization; its risk is whether developers receive usable permissions, debugging, distribution, and economics. The next test is whether the store opens on schedule and whether people complete tasks in real travel, work, and noisy environments.

SOURCES

Rokid official website

Yicai Global YodaOS report

Mogura VR Japan launch report

Rokid global product post


36Kr Rokid WAIC 2026 YodaOS launch

36Kr Rokid WAIC 2026 YodaOS launch

2026-07-02 to 2026-07-07 scan · 2026-07-11 follow-up · 2026-07-12 follow-up · 2026-07-13 follow-up · 2026-07-14 follow-up · 2026-07-14 follow-up · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

Z.ai ZCode packages GLM-5.2 as a China/global developer-tool signal

Product: Z.ai ZCode. Z.ai ZCode: An Agentic Development Environment for GLM-5.2. ZCode documentation says it brings GLM-5.2’s long-context, long-horizon, and agentic coding capabilities into a stable desktop experience for planning, coding, reviewing, and iterating across complex development tasks. Media coverage positions it against Cursor, GitHub Copilot, and Claude Code, with lower plan pricing reported. Prices, versions, and promotions should be checked against current Z.ai or ZCode official pages.. This item is handled as developer surface; specs, price, regions and dates use cited source claims only, and missing details remain source not stated.



Source-backed screenshot from ZCode's official docs, showing the developer product surface.

WHAT IT IS

An Agentic Development Environment for GLM-5.2. ZCode documentation says it brings GLM-5.2’s long-context, long-horizon, and agentic coding capabilities into a stable desktop experience for planning, coding, reviewing, and iterating across complex development tasks. Media coverage positions it against Cursor, GitHub Copilot, and Claude Code, with lower plan pricing reported. Prices, versions, and promotions should be checked against current Z.ai or ZCode official pages.

SPECS / STACK

The sources support ZCode, GLM-5.2, Agentic Development Environment positioning, macOS/Windows/Linux beta context, GLM Coding Plan, long-context and long-horizon work, planning/coding/reviewing/iterating, and account-plan quota metering. The IDE core, plugin protocol, repository permission model, full multi-agent architecture, enterprise SSO, code-privacy guarantees, regional restrictions, price stability, benchmark details, and safety evaluations are source not stated or require live official verification.

PAIN POINTS

ZCode addresses the gap between a capable coding model and an actual development workflow. A strong model is hard to use in real engineering if users must manually connect file context, permissions, long-running tasks, review steps, and quotas. ZCode packages those pieces into a desktop tool, so users do not have to assemble scripts around GLM-5.2 themselves. It also makes pricing, remote control, and multi-agent collaboration product-level choices.

NEW TECH

The product idea is an ADE rather than a small editor plugin. Model access, planning, execution, review, quota, remote bot control, and multi-agent coordination are combined into a development environment. It signals that Chinese AI companies are not only releasing models; they are competing for the developer workbench.

LIMITS / UNKNOWNNS

Open risks include large private-repository reliability, permission and secret handling, whether every change is explainable, support for local tests and rollback, whether UI similarity controversy affects trust, and whether GLM-5.2 performs consistently across English, Chinese, and mixed-language codebases.

HOW IT WORKS

A developer installs ZCode, connects a Z.ai or BigModel account and an active GLM Coding Plan, then starts a coding goal inside a project. The tool plans, edits code, reviews changes, and iterates across longer development tasks. Documentation and coverage also mention remote bot control through channels such as WeChat, Feishu, and Telegram, suggesting that a developer can trigger or supervise work from familiar communication tools instead of staying only inside an IDE.

USE CASES

The practical jobs are understanding large codebases with GLM-5.2, producing implementation plans, editing multiple files, reviewing changes, iterating on failures, remotely controlling agents from messaging tools, and trying a Chinese-model coding workflow at a lower reported price. For Chinese developers, it may reduce friction around local accounts, model access, payment, and communication channels. For global developers, it becomes a comparison point against Cursor, Claude Code, and Copilot.

USER VOICE

Media coverage notes online debate about low-cost competition and cloning accusations, but the cited sources do not establish representative user sentiment. Independent evidence on output quality, code safety, UI originality, long-term maintenance, and international experience is still thin; source not stated.

AVAILABILITY

ZCode documentation is available, and media reports describe a desktop application with Windows availability and macOS/Linux beta context. Actual downloads, pricing, promotions, platform support, account regions, and enterprise features should be verified through current Z.ai/ZCode documentation.

PRODUCT READ

ZCode is the most concrete China-lane developer-product signal today. It is not abstract model news; it is a model company entering the agentic IDE/ADE doorway. The user impact could be better pricing and better local ecosystem fit, but it must be validated on real repositories, test pass rates, permission design, and long-task recovery rather than launch-page claims.

GLOBAL

Global Signals

DEVELOPER SURFACE · DEVELOPER SURFACE · PLATFORM FACT
· PARTNER ROADMAP

Snapdragon Reality Elite pushes large-model capability into the XR platform

DEVELOPER SURFACE · DEVELOPER SURFACE · CONFIRMED
PARTNER PRODUCT

NVIDIA XR AI moves glasses from consumer entry point to field guidance

2026-06-16 official platform release · 2026-07 developer watch · 2026-07-14 follow-up · 2026-07-14 follow-up · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

Snapdragon Reality Elite pushes large-model capability into the XR platform

Product: Qualcomm Snapdragon Reality Elite. Qualcomm Snapdragon Reality Elite: Reality Elite is an XR platform for OEMs and developers, not a consumer glasses product sold on its own. Qualcomm’ s June 16 release positions it for all-in-one video-see-through headsets and lightweight tethered optical-see-through glasses, with language models, vision models, and spatial generative AI. It shows that low-latency agents on glasses are first a system problem of compute, graphics, tracking, and power. This item is handled as developer surface; missing details remain source not stated.

This website uses cookies from us and our business partners to enhance your browsing experience. For more information, including how to manage your cookie settings, please check the Privacy Policy at the bottom of this website and Qualcomm's [Cookie Policy](#).

1 Understand

Qualcomm official Reality Elite page screenshot; edge AI, display, and efficiency figures follow the official page.

WHAT IT IS Reality Elite is an XR platform for OEMs and developers, not a consumer glasses product sold on its own. Qualcomm’ s June 16 release positions it for all-in-one video-see-through headsets and lightweight tethered optical-see-through glasses, with language models, vision models, and spatial generative AI. It shows that low-latency agents on glasses are first a system problem of compute, graphics, tracking, and power.
SPECS / STACK Qualcomm lists up to 48 TOPS AI, up to 60% higher GPU, 30% higher CPU, and 160% higher NPU performance, with up to 4.4K per eye at 90 FPS. It claims up to 20% longer battery life at the same workload and a chipset up to 12°C cooler under load, compared with XR2+ Gen 2. The platform supports LLM/LVM, EVA vision acceleration, VST and OST, and comes first to XREAL Project Aura with a future Play for Dream device. Chip model, model size, SDK, price, and release date are source not stated.
PAIN POINTS Reality Elite attacks XR’ s system constraint: rich vision and spatial computing need power while glasses require low latency, low heat, low weight, and longer battery. Sending all visual understanding to the cloud adds network delay and privacy cost. A common NPU, GPU, CPU, display, and tracking platform gives OEMs headroom; the platform numbers still need a real device to prove model quality, heat, and battery can coexist.
NEW TECH The stack combines LLM and LVM support, Gaussian Splatting, EVA vision acceleration, head and hand tracking, and a see-through pipeline in scalable XR silicon. It is not a complete on-device agent, but it makes local perception and inference more practical. The HCI consequence is that teams must decide which feedback can arrive within interaction time, which waits for the cloud, and how users understand that difference.
LIMITS / UNKNOWNNS Open questions include sustained power at 48 TOPS, thermal behavior in glasses, model memory use, offline behavior, camera and microphone data boundaries, and whether OEMs turn platform capability into legible permissions and state. Project Aura’ s final specifications, price, timing, and consumer experience remain roadmap or partner signals and cannot be inferred from the chip release.

HOW IT WORKS The platform supplies real-time vision, head and hand tracking, see-through processing, and on-device AI rather than one consumer interface. OEMs can send voice, gesture, gaze, and environmental vision into LLM or LVM services and return object recognition, spatial content, or an agent response through lens and audio layers. Qualcomm cites object generation and real-time spatial awareness; wake, permission, error, and stop behavior remain OEM choices.
USE CASES The platform supports visual generation, spatial measurement, object recognition, tracking, immersive content, and agent guidance. Potential jobs include understanding nearby objects, generating content in space, receiving environmental cues, and connecting field data to a digital twin in industrial or training work. Public sources confirm the platform and partner roadmap; they do not prove a consumer product has delivered a stable all-day agent at 48 TOPS.
USER VOICE Source not stated.
AVAILABILITY Qualcomm’ s official product and release pages are public. Qualcomm says Reality Elite will first underpin XREAL Project Aura and a future Play for Dream device. Developer resources, partner hardware, SDK access, retail launch, regions, and end-product price are not fully disclosed. The platform must not be described as a consumer product already available for purchase.
PRODUCT READ Reality Elite is infrastructure evidence for agent-first XR, not evidence that a finished product exists. It turns ‘can glasses understand, respond, and remain wearable?’ into a joint problem of silicon, power, optics, and tracking. Edge AI will determine latency, privacy, and failure modes; the proof must come from partner devices, not a TOPS number.

Global developer surface · developer surface · confirmed partner product

2026-06-16 developer release · 2026-06 Helix product reveal · 2026-07-13 follow-up · 2026-07-14 follow-up · 2026-07-14 follow-up · 2026-07-16 follow-up · 2026-07-20 follow-up · 2026-07-21 follow-up

NVIDIA XR AI moves glasses from consumer entry point to field guidance

Product: NVIDIA XR AI / VITURE Helix. NVIDIA XR AI / VITURE Helix: NVIDIA XR AI is a developer library and platform for connecting real-time vision, voice interaction, and agent guidance to AR glasses. VITURE Helix is an AI safety-glasses product built with that solution. NVIDIA’ s examples include an engineer asking an agent about a programmable-logic-controller issue through glasses while connecting industrial systems, digital twins, and automation workflows. The developer page also shows laboratory and gene-editing scenarios. Unlike a consumer glasses pitch, the value is not only capture or Q&A; it is putting the next step, procedure, and equipment context into the worker’ s view and voice loop.. This item is handled as developer surface; specs, price, regions and dates use cited source claims only, and missing details remain source not stated.



Source-backed screenshot from the NVIDIA XR AI developer page showing the developer surface for real-time, hands-free agent guidance.

WHAT IT IS

NVIDIA XR AI is a developer library and platform for connecting real-time vision, voice interaction, and agent guidance to AR glasses. VITURE Helix is an AI safety-glasses product built with that solution. NVIDIA’ s examples include an engineer asking an agent about a programmable-logic-controller issue through glasses while connecting industrial systems, digital twins, and automation workflows. The developer page also shows laboratory and gene-editing scenarios. Unlike a consumer glasses pitch, the value is not only capture or Q&A; it is putting the next step, procedure, and equipment context into the worker’ s view and voice loop.

SPECS / STACK

The cited sources support the NVIDIA XR AI developer library, LabOS, compatibility with Meta, Rokid, and VITURE glasses, and VITURE Helix’ s partnership with NVIDIA XR AI. The developer page emphasizes real-time, hands-free guidance and examples involving industrial PLCs, laboratories, and gene editing. VITURE supplies the smart hardware and edge software. The sources do not specify Helix’ s chip, camera parameters, display specifications, weight, battery, network protocol, model name, edge/cloud split, SDK version, price, or purchase regions; those details are source not stated.

PAIN POINTS

The product attacks the friction of knowing what to do while being unable to leave the task. A field worker traditionally stops, picks up a phone or tablet, searches documentation, confirms the model, and returns attention to the machine. XR AI attempts to compress that loop into look, ask, hear, and act. The cost is higher consequence for wrong guidance. The visual agent must handle occlusion, lighting, equipment variation, and step state. A professional product needs citations, confidence, human takeover, and stop controls rather than only a smooth demo.

NEW TECH

The new technology pattern treats XR hardware as an agent’ s first-person sensor and feedback endpoint, not merely as a remote display. NVIDIA provides an XR AI library; LabOS connects visual input, domain tools, and agent orchestration; VITURE brings the pattern into safety-glasses hardware. The HCI challenge moves from 'where should the UI sit?' to 'how does field context enter the model, and how does the model produce an executable but auditable next step?'

LIMITS / UNKNOWNNS

Open questions include field connectivity, edge latency, occlusion and lighting robustness, knowledge integration for different factories, hallucination, compliance logging, operator privacy, multi-user coordination, wear fatigue, and responsibility after an error. If safety glasses provide advice, the product must separate advice from permission to execute. If they later control automation, permissions and two-person confirmation become hard requirements.

HOW IT WORKS

A worker wears compatible glasses in the field. The glasses or companion compute captures first-person visual context and the user’ s spoken question. XR AI components help an agent interpret equipment, work steps, or experiment context, then return an explanation, reminder, next-step instruction, or remote-collaboration signal through audio or an in-view display. VITURE positions Helix as an eyes-on, hands-free AI co-pilot; LabOS examples connect perception to digital twins, automation, or experiment workflows. A high-risk action still requires professional SOP and human authorization; a guidance demo is not automatic control permission.

USE CASES

Concrete jobs include maintenance workers receiving procedure prompts while looking at equipment, engineers diagnosing PLC or automation systems, lab staff retrieving context at the bench, trainers guiding a new operator remotely, quality inspection, and any field task where both hands need to stay on the work. For enterprises, glasses do not need to replace a computer. They can place computer vision, a knowledge base, digital twins, and agent recommendations where the work actually occurs. The product should be judged by time saved on manuals, workstation switching, and repeated confirmation without creating a larger safety risk.

USER VOICE

Current evidence is NVIDIA and VITURE product material plus demonstration context, not enough independent long-term evidence from factories or labs. Early VITURE community discussion raises questions about Android support, product positioning, and whether real-time SOP guidance works outside the demo. Those are early friction signals, not validated performance data; source not stated.

AVAILABILITY

The NVIDIA XR AI developer page and blog are public. VITURE’ s Helix product page and launch article are live and identify the product as an AWE 2026 collaboration. Compatible glasses, developer access, purchase channels, customer industries, delivery timing, service pricing, and regional supply are not fully disclosed by the current sources; source not stated.

PRODUCT READ

NVIDIA XR AI and Helix are the strongest field-agent signal in this issue because they move glasses from a lightweight consumer entry point into a professional work system. The opportunity is real only if the system produces sourced, pausable, handoff-ready next steps in the field. For HCI teams, display area is the surface problem; context, responsibility, confidence, and human takeover are the system problem.

DESIGN DESK

Design Desk: Design the state that earns continuous wear

Today's evidence compresses the glasses product into a more concrete chain: whether weight earns wear, whether capture earns repeated use, whether display deserves the field of view, whether camera state is verifiable to bystanders, whether an agent can report while the user moves, how failure stops, and how data is deleted.

01

Weight must become habit

MOONIX targets approximately 14.9 grams while RayNeo spends 76 grams on binocular display; both routes require content that is short, frequent, and worth keeping on the face.

02

Camera boundaries must be verifiable

Meta couples a capture LED to camera disablement on tamper; MemoMind removes the camera. Public trust needs visible state, user verification, and enforced behavior.

03

A display is not a shrunk phone

Binocular MicroLED, optical see-through, and Compose Glimmer demand glanceable hierarchy; every cue needs a stop, return, and recovery path.

04

Agent state must travel off-screen

Monako imagines coding-agent waiting, error, and confirmation as HUD states; community discussion asks glasses agents to create apps, navigate, and run timers, so execution needs permissions, tool visibility, state change, and undo.

05

Edge inference must explain waiting

OpenGlass points to event-driven power budgets; products need to show when work stays local, when data leaves, and when the system degrades.

06

Captured data must be deletable

MOONIX capture, MemoMind voice notes, OpenGlass/VisionClaw sensing, and patent material generate validation questions; upload, retention, export, and deletion must become visible states.

NEXT SCAN

What to watch next

01

MOONIX: hands-on runtime for 14.9g, 16 hours, and six microphones under continuous mixed capture; recording state, deletion, and disconnected fallback on the standard model.

02

Monako Glass: July-August delivery, whether Lua apps really run without a build step, agent permission and secret isolation, and HUD legibility.

03

Meta/Android XR: model coverage for LED tamper safeguards, false-positive recovery, daylight bystander visibility, and public-space cues.

04

Google/Samsung Intelligent Eyewear: actual autumn hardware, permission confirmation, camera state, and disconnect fallback.

05

Rokid YodaOS: real devices after the WAIC launch, ecosystem APIs, model routing, and channel delivery.

06

RayNeo X3 Pro: mixed-load runtime, long-term comfort at 76 grams, Creator Mode, and regional Gemini behavior.

07

MemoMind and iFlytek: delivery, camera-free privacy, language-by-language quality across 122 languages, and GlassClaw send confirmation.

08

Community agent use, OpenGlass/VisionClaw, and smart-glasses patents: keep community, research, and patent labels until reproducible product evidence appears.

Every topic has inline sources; this is the quick ledger.

SOURCE MOONIX official product page	SOURCE MOONIX official English product page	SOURCE MOONIX downloads / companion app	SOURCE IT之家 launch and availability report	SOURCE 深圳湾 product interview	SOURCE Monako official website
SOURCE Product Hunt Monako Glass listing	SOURCE Product Hunt launch discussion	SOURCE Buildroot official downloads	SOURCE RayNeo X3 Pro official UK product page	SOURCE TechRadar Pro hands-on review	SOURCE RayNeo lens partner Lensology
SOURCE MemoMind One official product page	SOURCE Android Central hands-on review	SOURCE Tom’ s Guide hands-on review	SOURCE MemoMind Kickstarter / pre-order	SOURCE 36Kr iFlytek AI Glasses launch report	SOURCE iFlytek 2025 annual report PDF
SOURCE IT Home launch report	SOURCE iFlytek official website	SOURCE Google intelligent eyewear preview	SOURCE Android XR Developer Preview 4	SOURCE Android XR Gemini API guide	SOURCE Android XR platform
SOURCE Android Central anti-tampering interview	SOURCE Meta AI Glasses privacy FAQ	SOURCE Meta glasses privacy approach	SOURCE Android Central hands-on	SOURCE Rokid official website	SOURCE Yicai Global YodaOS report
SOURCE Mogura VR Japan launch report	SOURCE Rokid global product post	SOURCE 36Kr Rokid WAIC 2026 YodaOS launch	SOURCE HTC VIVE Eagle official launch	SOURCE HTC VIVE Eagle product page	SOURCE HardwareZone hands-on review
SOURCE HTC security and privacy whitepaper	SOURCE Brilliant Labs Halo product page	SOURCE Brilliant Labs developer page	SOURCE Halo developer documentation	SOURCE Brilliant Labs and Alif partnership	SOURCE Brilliant community shipping discussion
SOURCE Snap official SPECS launch	SOURCE Snap SPECS product page	SOURCE TechCrunch SPECS launch report	SOURCE Snap Lens Studio developer tools	SOURCE Even G2 official product page	SOURCE TechCrunch Even G2 hands-on
SOURCE Even Hub developer surface	SOURCE Even support surface	SOURCE Qualcomm Reality Elite release	SOURCE Qualcomm Reality Elite product page	SOURCE Qualcomm product brief	SOURCE XREAL Project Aura
SOURCE Microsoft MDEP: Project Solara	SOURCE Project Solara coverage	SOURCE Microsoft Entra	SOURCE Microsoft Intune	SOURCE Reddit r/augmentedreality agent-use post	SOURCE Reddit r/augmentedreality WAIC post
SOURCE Reddit smart-glasses navigation test	SOURCE OpenAI: Introducing GPT-Live	SOURCE OpenAI live replay	SOURCE MacRumors: GPT-Live voice coverage	SOURCE TechRadar: GPT-Live rollout	SOURCE NVIDIA Developer: XR AI Platform
SOURCE NVIDIA Blog: XR AI brings agents to AR glasses	SOURCE VITURE: Helix with NVIDIA XR AI	SOURCE VITURE Helix product page	SOURCE ZCode docs: welcome	SOURCE ZCode docs: configuration	SOURCE Business Insider: Z.ai ZCode launch
SOURCE DevOps.com: Z.ai debuts ZCode	SOURCE OpenGlass arXiv	SOURCE OpenGlass HTML	SOURCE Prophesee GENX320	SOURCE VisionClaw arXiv research signal	SOURCE WIPO Global Awards 2026: Rokid
SOURCE Google Patents: AI-supported smart glasses	SOURCE Google Patents: AI Smart Glasses GS1 design	SOURCE Android Central: Solos smart-glasses patent lawsuit	SOURCE Google Patents conversational smart-glasses assistant		

Full ledger: [sources.md](#)