

AI Daily | New products and useful questions, about 15 minutes

2026-10-06 • America/Toronto • 2026-10-06T16:17:28-04:00

Contents

12 dated recent advances and 4 first-inclusion background stories with original publication dates; more physical AI hardware, real workflows and demonstrations on integrated story pages

01 Ghost Core

2026-10-05

02 Coco

2026-10-05

03 OpenVision

2026-10-04

04 Mulligan

2026-10-05

05 Beam

2026-10-05

06 BazaarBench

2026-10-05

07 HyperFrames Studio

2026-10-06

08 Cua Driver

2026-10-05

09 LearnSched

2026-10-05

10 SafeWorld

2026-10-05

11 North 2

2026-10-05

12 Claude Cowork

2026-10-06

13 Surface

Background • not new today • 2026-09-24

14 Omega AI+

Background • not new today • 2026-10-01

15 Digit 5

Background • not new today • 2026-10-01

16 Patent context

Background • not new today • 2026-10-01


Ghost Core packages local agents as a dedicated computer



Ghost's official front product render [\[1\]](#)



Ghost's official rear view: ventilation and power input [\[1\]](#)



First frame of the official hardware animation: the enclosure is still assembled [\[1\]](#)

What changed

Preorders opened October 5. Check the official page for current ordering availability and delivery dates; this is not an independent delivered-unit review [\[2\]](#)

What it is

A screenless US\$3,499 personal AI computer marketed for local models and memory

When to use it

It bundles a 24GB RTX PRO 4000 Blackwell GPU, 64GB RAM, 1TB SSD and preinstalled open-weight models

Limits & availability

Continuous context and privacy are vendor claims, not independent tests. Local inference doesn't establish that external-app communications stay offline

What the box contains


The CPU is Ryzen 5 7600; the listed models are Qwen 3.8-Next, Qwen 3.8-27B, Gemma 4-31B and Muse-Glimmer-30B [\[1\]](#)

The page groups context into digital, home and health; these illustrate product intent rather than prove every integration is mature [\[1\]](#)

Play original demo • Official launch-video link [\[38\]](#) [\[1\]](#)

Product form and interaction presentation; the external video was inaccessible, so no unverified timestamps are supplied

Coco announces UK robot delivery and pavement-hazard sharing



Official UK partnership-page image; not evidence of nationwide operation [3]

What it is

Electric sidewalk robots carry orders to customers

What changed

A UK deployment plan starts in Canary Wharf in October, followed by four cities; local merchants and charging depots form part of the rollout [3]

When to use it


Eligible customers order food and receive robot-carried deliveries. Onboard cameras, sensors and edge AI perceive the route; hazard observations will also feed BlindSquare's spoken navigation

Limits & availability

Phased service zones, no exact UK opening day, and most orders still use riders. Local speed/routing rules and supervising operators remain important; edge navigation is not a claim that the whole service runs offline



Coco official photograph: delivery robot on snow; not proof of nationwide UK availability [39]



A third-party open-source iPhone app makes glasses a first-person interface to chosen local or cloud models

Current main-branch fixes improve credential persistence and failed-speech recovery; this is neither Meta firmware nor an App Store launch [5]


Build and register the app, connect glasses, choose a backend, then ask about what you see and hear its reply. October 2 added subscription sign-in, with capability differences across backends

Source-build and developer setup required. Subscription sign-in uses an undocumented CLI surface and may disrupt another session; live video still needs an API key or another backend. Failed Keychain writes can retain a secret in the old file


Play original demo • Earlier developer demo: a seen object leads to browser search [41] [40]

Observe the Coke can and browser product search; this is not a test of current fixes

Mulligan targets human-supervised robot practice at observed failures



Author demo at 12 seconds: one policy, two initial states; one succeeds and one fails to grasp [11]



Author demo at 24 seconds: evaluate evenly across the full initial-state range [11]

What it is

A Stanford human-in-the-loop research system runs on a Franka Panda/DROID arm, not a new retail robot

What changed

Three physical tasks and 2,550 held-out blinded episodes test a method that selects where objects should start, rather than repeating easy placements [10]

When to use it

An operator follows reset guidance, watches insertion or cable routing, takes over via SpaceMouse when needed, and retrains between rounds


Limits & availability

The 98% collection-completion figure includes human interventions. Reported 10–34-point gains concern selected precision tasks, not general household autonomy

Play original demo • Author walkthrough of a Thread nut collection and improvement round [42] [11]

See failed placements feed the next round, with human resets and interventions kept explicit

The next round starts at the 9 failures, plus new initial states.




Author demo at 40 seconds: retry nine failed starts and add underexplored placements [11]

Beam preview: distinguish active parameters from deployment cost



Official reasoning-effort figure: generated tokens versus task score [13]



First frame of the official NYC subway-map demo [13]

What it is

A 501B-total, 23B-active MoE aimed at coding and agent workloads

What changed

Early-access signup opened October 5; Apache 2.0 weights and the report remain promised for later this month [13]

When to use it

An official demo shows a generated NYC subway map; plots compare reasoning effort, scores and token use


Limits & availability

Vendor compute estimates omit prefill, attention and serving overhead, so they aren't measured user costs

Play original demo • Official demonstration of a subway map created with Beam [43] [13]

Watch the map zoom and line/train markers change; this is a vendor-selected output example

BazaarBench asks why completed trades can still be unsafe



Paper Figure 1: market, inventory records, agent actions and safety checks [15]

What it is

A new King's College London/Oxford-linked preprint tests agents trading for users

What changed

Across five models and 45 simulated continuations, all attempted double commitments under ordinary instructions [15]

When to use it

Check inventory, condition and prior promises alongside task completion

Limits & availability


A simulated marketplace, not a real-world fraud rate; some labels use an LLM judge. Its long-term influence is not established

Read the experiment carefully

Each base market has 100 agents; continuations preserve histories while changing the same 20 agents' model and instructions [15]

Accepting an offer doesn't automatically reserve stock, so avoiding conflicting promises is the agent's responsibility [15]

HyperFrames Studio shares the timeline; today's patch improves preview memory



Official Home screenshot: begin with a prompt, with projects listed on the left [\[18\]](#)

What it is

HeyGen's desktop editor lets a coding agent and a person work on one video project

What changed

Today's v0.8.137 reduces preview-loading memory for large inline assets and flags excessive audio peaks; desktop setup guides arrived October 5 [\[16\]](#)


When to use it

Describe a cut, then point, draw, comment on a frame or adjust the timeline before sending grouped edits to the agent

Limits & availability

Official setup covers macOS and Linux x86_64, with HeyGen sign-in and your own agent; a Windows download wasn't verified

Cua Driver makes agent cursor motion easier to read



Official session-cursor and badge examples show targets and delivery state; this still image does not demonstrate the six motion trajectories [21]

What it is

A visible cursor overlay for agents operating desktops

What changed

Six styles include signature arc, spring settle, magnetic, comet swoop, adaptive and classic [20]


When to use it

Motion, click feedback and session labels can clarify the action being observed

Limits & availability

A visual aid, not an authorization credential. Reduced-motion settings override effects; improved task accuracy wasn't independently established

LearnSched: idle GPUs need not immediately self-improve



LearnSched preprint Figure 1: a shared GPU trades current service against future improvement and pays a recovery cost [22]

What it is

Yangbo Wei and eight co-authors study GPU sharing between agent service and self-improvement

What changed

LearnSched accounts for progress, checkpoints and recovery across 1,920 paired finite-model scenarios [22]

When to use it

Estimate the reuse window before starting or pausing background improvement

Limits & availability

Not a production deployment guarantee; a simpler window policy matches value in evaluated warm-recovery cases

SafeWorld funds the testing layer before robots meet people



SafeWorld's official warehouse human-robot simulation illustration; not a test result, certification or real-world incident-rate evidence [24]

What it is

Robot safety-simulation tools for environments shared with people

What changed

A US\$12.2M seed round co-led by Shine Capital and a16z Speedrun was announced October 5 [23]


When to use it

Build browser scenarios from incidents, standards and robot logs, then repeat them with varied human motion


Limits & availability

The company describes early pilots and request-access availability; simulation coverage is not certification or a guarantee of zero physical incidents

North 2 brings reusable agents, workflows and spending limits into one workspace



Official interface illustration: connected knowledge and tools; figures are demo content [44]



Official interface illustration: document creation and export; figures are demo content [44]

What it is

Cohere's North 2 is an enterprise agent workspace connected to company knowledge and tools

What changed


The October 5 upgrade adds rebuilt orchestration, shareable agents, memory, skills, libraries and finer quotas/spending controls [25]

When to use it

Build automations from prompts and visual workflows, or create documents and lightweight apps in chat


Limits & availability

Access is sales/demo-led; pricing and customer-specific activation dates aren't stated. Some financial-data connectors remain planned



Official interface illustration: recurring automation setup; figures are demo content [44]

Claude Cowork moves new Pro/Max tasks to the cloud today



Official product demo: task beside a built-in browser. Not a post-change hands-on capture; the October 6 cloud change is evidenced by the Help Center [27]

What changed

New tasks run in the cloud and the local-only option disappears; already-started local tasks can continue [26]

What it is

Claude's long-running task workflow changes execution location for Pro/Max today

When to use it

Cloud work can continue with a closed laptop and be reviewed across desktop, web and mobile; schedules without local dependencies need no online device


Limits & availability

Local files still require an open desktop app and connected folders. Local projects and schedules don't transfer directly to Claude Code

Surface adds pen-first AI context and haptic mouse feedback



Microsoft official product image: new Surface Pro, magnetic pen and Surface Mouse [28]



Microsoft official Ink Canvas screen: handwriting and image generation on one canvas [28]

What it is

Physical input products pair a Surface Pro 12-inch tablet and Slim Pen AI canvas with a new ambidextrous Surface Mouse

What changed

Announced September 24, with selected-market availability October 13. Mouse starts at US\$79.99; computer, pen and AI plan are separate costs [28]

When to use it

Draw, write and annotate on Ink Canvas before generating imagery or summarizing text. Customize the mouse action button for Copilot or shortcuts; supported clicks, window snapping and alignment give haptic feedback

Limits & availability

Pen and keyboard are separate. Canvas AI requires an eligible Microsoft 365 plan and has usage limits. Mouse customization needs the Windows Surface app; haptics are ordinary UI feedback, not proof an agent task succeeded



Microsoft official Surface Mouse image; action-button and haptic feedback are the input changes [28]

Omega AI+ puts adaptive sound and broadcast controls at the ear



Manufacturer launch image: Omega AI+ hearing aids and charger; not evidence of clinical benefit [30]

What it is

A new hearing-aid family runs speech/noise processing at the ear, with visible phone controls for adjustments and connections

What changed

Launched October 1; the October 5 official explainer adds workflow detail. Earlier Omega AI performance is not AI+ evidence [30]

When to use it

After individualized fitting, automatic adaptation handles changing sound. My Starkey lets wearers change programs, volume or Edge Mode, select compatible Auracast broadcasts, and receive the new left-behind reminder

Limits & availability

Benefits are manufacturer claims, with no independent clinical verification here. Wireless/custom styles differ; sleep features vary by location. No universal price or market-delivery table is given; fitting and suitability need professional evaluation




Official usage illustration: wearer views My Starkey hearing controls; promotional scene, not an independent review [31]


Digit 5 extends safety from the robot to the surrounding facility



Agility official Digit 5 scene: tote and workstation; vendor product imagery, not a new field test [35]



Agility official fleet-interface illustration: robot ID, battery, docking and Auto state [35]



Agility official scene: robot seated with a worker nearby; imagery is not independent safety validation [35]

What it is

A warehouse/factory humanoid product; the current development is its deployment safety layer, not home retail

What changed

The planned offboard bridge lets surrounding safety systems participate. October 1 is a memorandum and planned architecture, not proof every component is deployed [33]

When to use it

Normal work runs autonomously; independent human detection can trigger avoiding, stopping or sitting, while operators monitor and override through a tablet pendant with emergency stop and enabling-device functions

Limits & availability

Digit 4' s 65,000 hours are not Digit 5 performance. No universal delivery date or retail price is supplied; vendor safety/productivity claims need certification and site-specific checks

Patent context: why wrist gestures need remapping after repositioning




Figure 1A

Figure 1A: initiate micromovement guidance [36]




Figure 1E

Figure 1E: try moving the little finger [36]

What it is

Meta Platforms Technologies' biopotential gesture/feedback application
US20260299703A1

What changed

Published October 1; filed March 16, 2026; earliest claimed provisional priority October 19, 2023 [36]

When to use it

Figures 1A–1M illustrate placement, little-finger practice and progressively smaller movements

Limits & availability

Publication is neither a grant nor shipping evidence; descriptive examples and current independent claims must be read separately

Read drawings and claims separately

This is a continuation of an October 20, 2024 application and claims a 2023 provisional. Separate dates prevent years of work from being presented as today's invention [36]

Claim 1 is cancelled. Independent claims 2,9,16 include initial gesture mapping, repositioning detection, a new representation and later recognition using it [36]

The coaching figures are descriptive examples; this report makes no novelty, patentability or infringement determination [36]




Figure 1G

Figure 1G: minimize little-finger movement [36]

References & evidence index

Full dates, figure provenance & selection notes → [sources.md](#)

n.d. = publication date not stated; effective dates are labeled separately

- | | | |
|--|--|-------------------------------------|
| 01 Ghost • product • n.d. | 16 HyperFrames • v0.8.137 • 2026-10-06 | 31 Starkey • product • 2026-10-05 |
| 02 TechCrunch • article • 2026-10-05 | 17 HyperFrames • v0.8.134 • 2026-10-05 | 32 Starkey • product • n.d. |
| 03 Coco • release • 2026-10-05 | 18 HyperFrames • guide • n.d. | 33 Agility • article • 2026-10-01 |
| 04 Deliveroo • release • 2026-10-05 | 19 HyperFrames • guide • n.d. | 34 Agility • release • 2026-09-15 |
| 05 GitHub • article • Effective date: 2026-10-04 | 20 Cua • release • 2026-10-05T22:30:00Z | 35 Agility • guide • n.d. |
| 06 GitHub • article • 2026-10-04 | 21 Cua • guide • n.d. | 36 USPTO • patent • 2026-10-01 |
| 07 GitHub • article • 2026-10-04 | 22 arXiv • 2610.06212 • 2026-10-05T12:22:40Z | 37 Patentlyze • patent • 2026-10-02 |
| 08 GitHub • v2.14.0 • 2026-10-02 | 23 WebWire • release • 2026-10-05 | 38 X • demo • n.d. |
| 09 GitHub • article • n.d. | 24 SafeWorld • product • n.d. | 39 Coco • figure • n.d. |
| 10 arXiv • 2610.05882 • 2026-10-05 | 25 Cohere • article • 2026-10-05 | 40 GitHub • article • n.d. |
| 11 Mulligan • demo • n.d. | 26 Claude • article • Effective date: 2026-10-06 | 41 Source • GIF • n.d. |
| 12 arXiv • 2610.05882v1 • n.d. | 27 claude.com • demo • n.d. | 42 Amazon S3 • video • n.d. |
| 13 Reflection • article • 2026-10-05 | 28 Microsoft • release • 2026-09-24 | 43 Sanity • GIF • n.d. |
| 14 arXiv • 2610.06748 • 2026-10-05T17:26:17Z | 29 Microsoft • article • 2026-09-24 | 44 Cohere • demo • n.d. |
| 15 arXiv • 2610.06748v1 • 2026-10-05 | 30 PRNewswire • release • 2026-10-01 | |